

2IC30: Computer systems Instruction Execution: Speeding Up by Prefetching. Processor basics: Binary and Assembly Language; ISA instruction types.

Jan Friso Groote

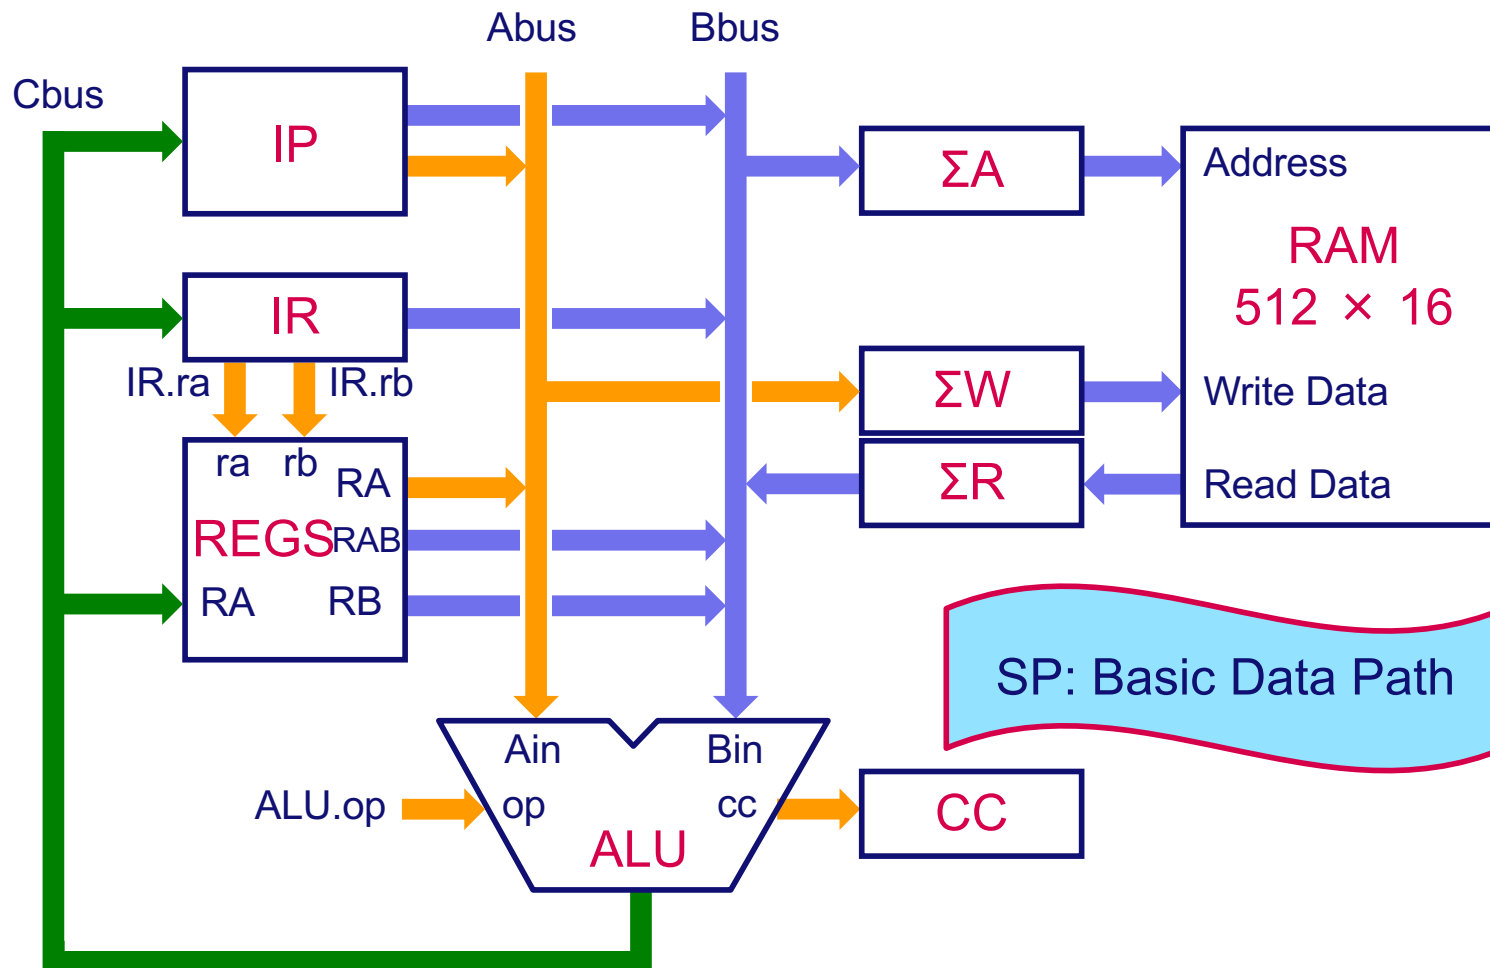


TU / **e**

Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

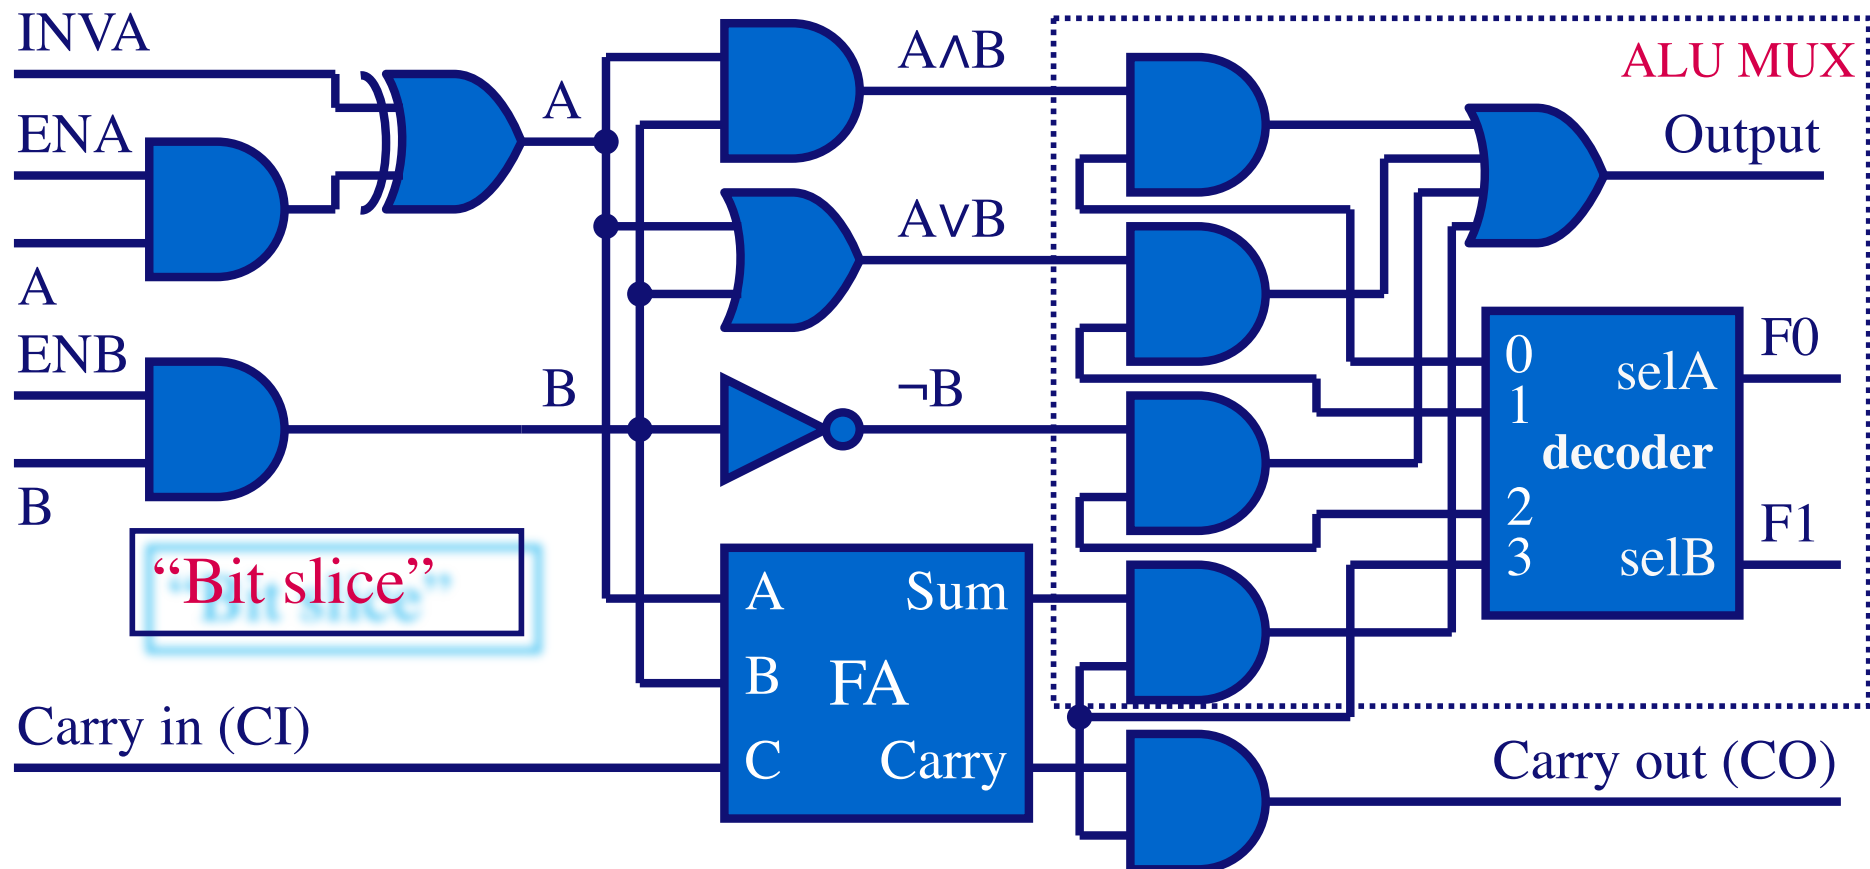
Simple processor: data path



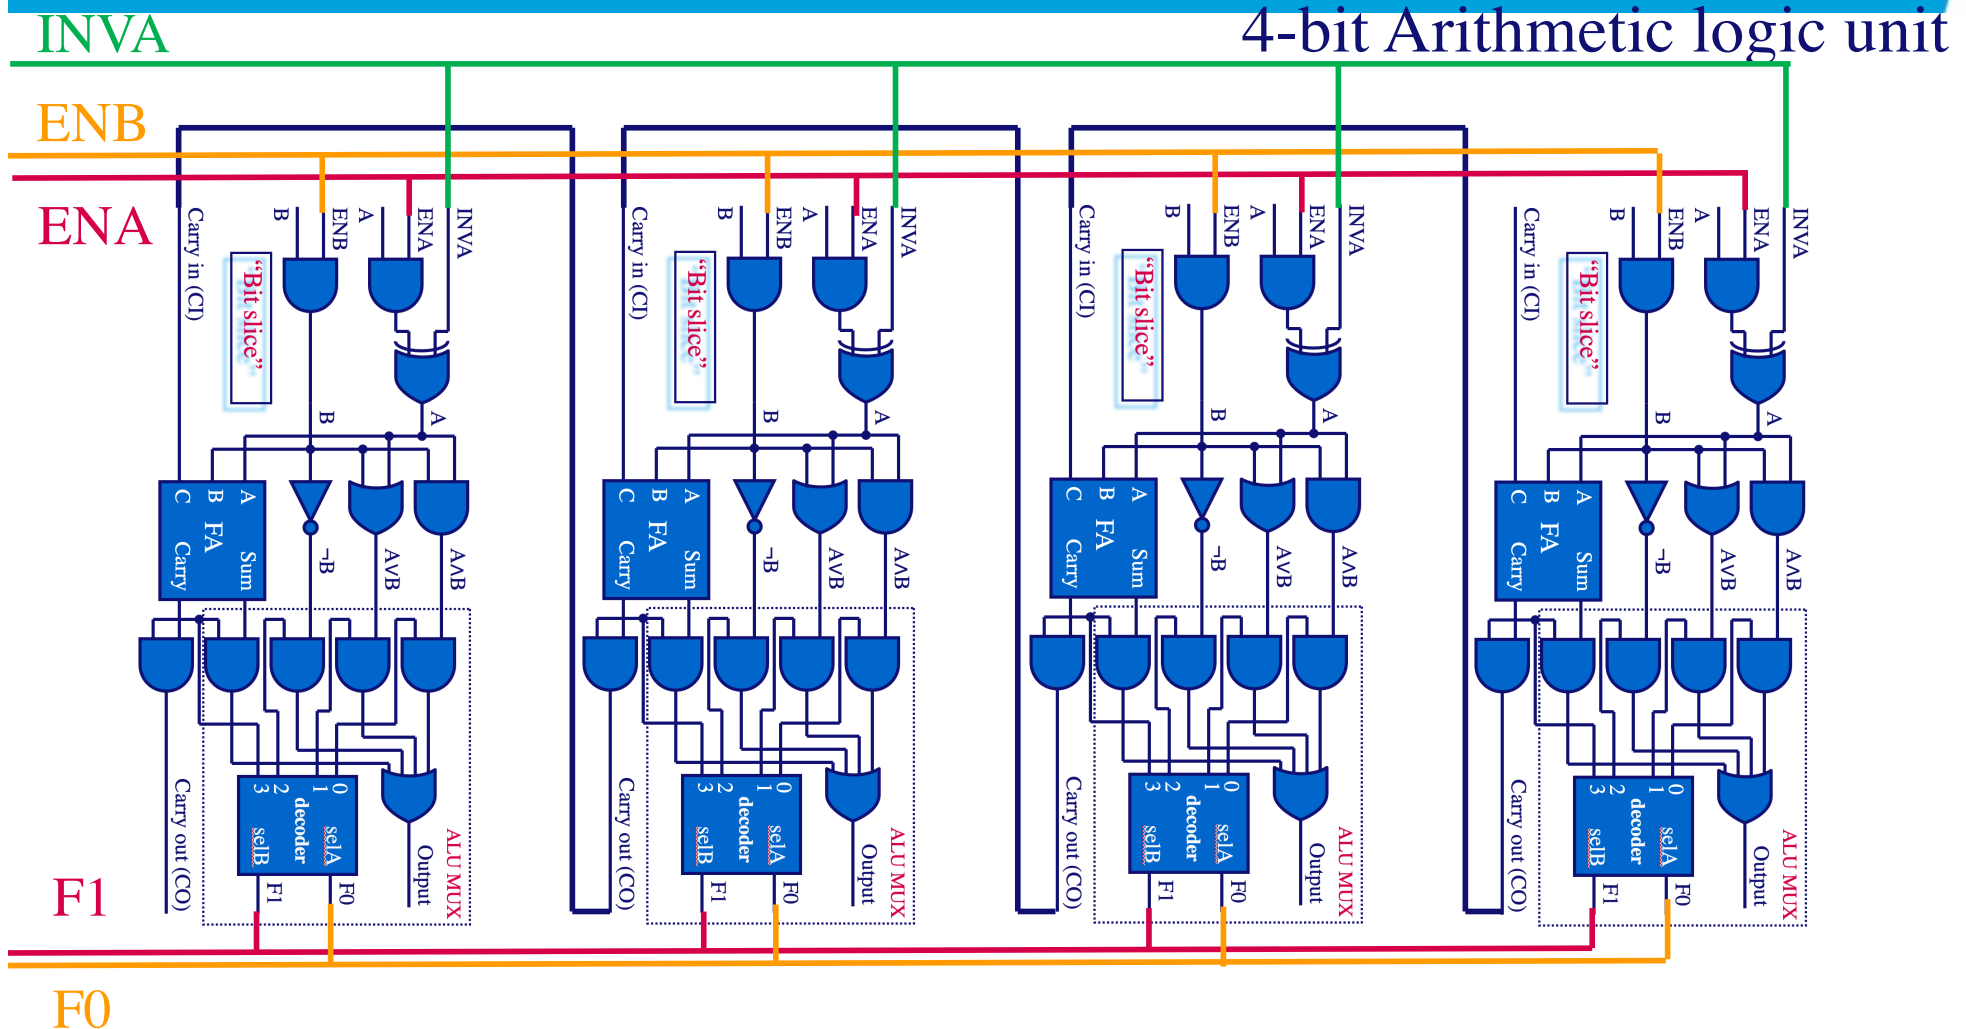
- Abus: IP or RA
- Bbus: IP, IR, ΣR, RB or RAB
- REGS: 2 Reads and 1 Write, simultaneously, ra[1..0] and rb[1..0] select registers

The arithmetic logical unit (1)

- ❑ Hardware for one bit is doable (the ALU in the reader is different)...




The arithmetic logical unit (2)



- ❑ Connect “bit slices” through carry in / carry out
- ❑ Control with F0, F1, ENA, ENB and INVA (equal for all slices)
- ❑ Tests: add logic for Negative, oVerflow and Zero

The arithmetic logical unit (2)

- ❑ Connect “bit slices” through carry in / carry out
- ❑ Control with F0, F1, ENA, ENB and INVA (equal for all slices)
- ❑ Tests: add logic for Negative, oVerflow and Zero



Functions of the example ALU

- ❑ Five control signals and C_i : $2^6 = 64$ functions ?

- Well yes, but some of those are useless...

- ❑ Pass values and generate constants:

$A, B, 0$ (zero), 1 (one), -1 (minus one)


- ❑ Mathematical functions:

$A+B, A+B+1, B-A, B-A-1,$
 $0-A$ (NEG), $A+1$ (INC), $B+1, B-1$ (DEC)

- ❑ Logical functions:


$A \text{ AND } B, A \text{ OR } B, \text{NOT } A, \text{NOT } B$

The ALU from the book.



The ALU in the book
is slightly different.

The 4-bit ALU from the book.



Simple processor: ALU functions

ALU functions for our simple processor (see the reader)

loadA: A

loadB: B

inc: $A + 1$

dec: $A - 1$ (not with bitslice ALU)

add: $A + B$

sub: $A - B$ (not with bitslice ALU)

and: $A \wedge B$

or: $A \vee B$

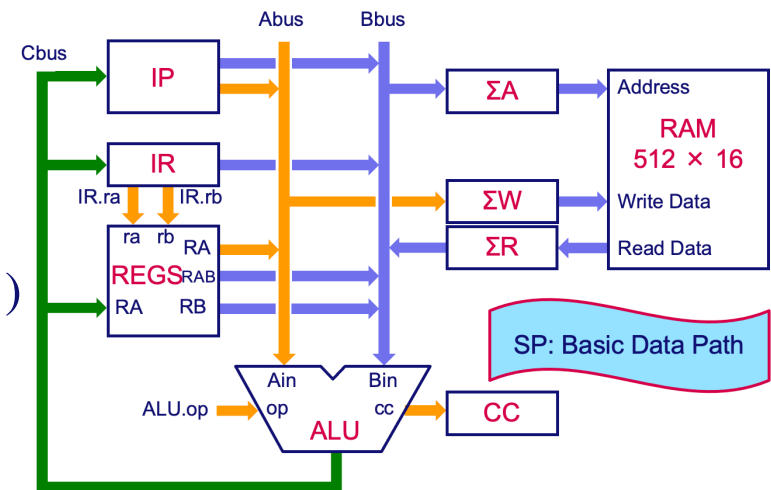
xor: $A \oplus B$ (not with bitslice ALU)

Conductor: instruction execution (ADD, direct)

Instruction: ADD RA RB

Required action: $RA \leftarrow RA + RB$

- 0 $\Sigma A, IP \leftarrow IP(Bbus), IP(Abus)+1$ (ALU: “A+1”)
- 1 $\Sigma R \leftarrow RAM[\Sigma A]$
- 2 $IR \leftarrow \Sigma R$ (ALU: “B”)
- 3 $RA, CC \leftarrow RA + RB$ (ALU: “A+B”), ALU.cc



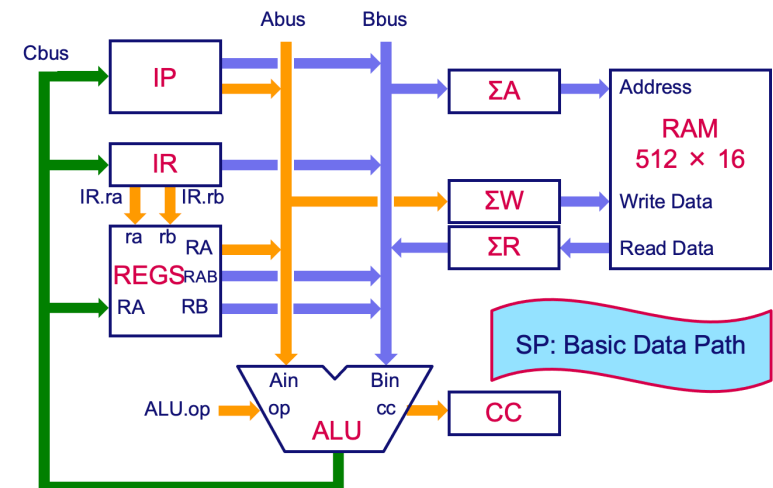
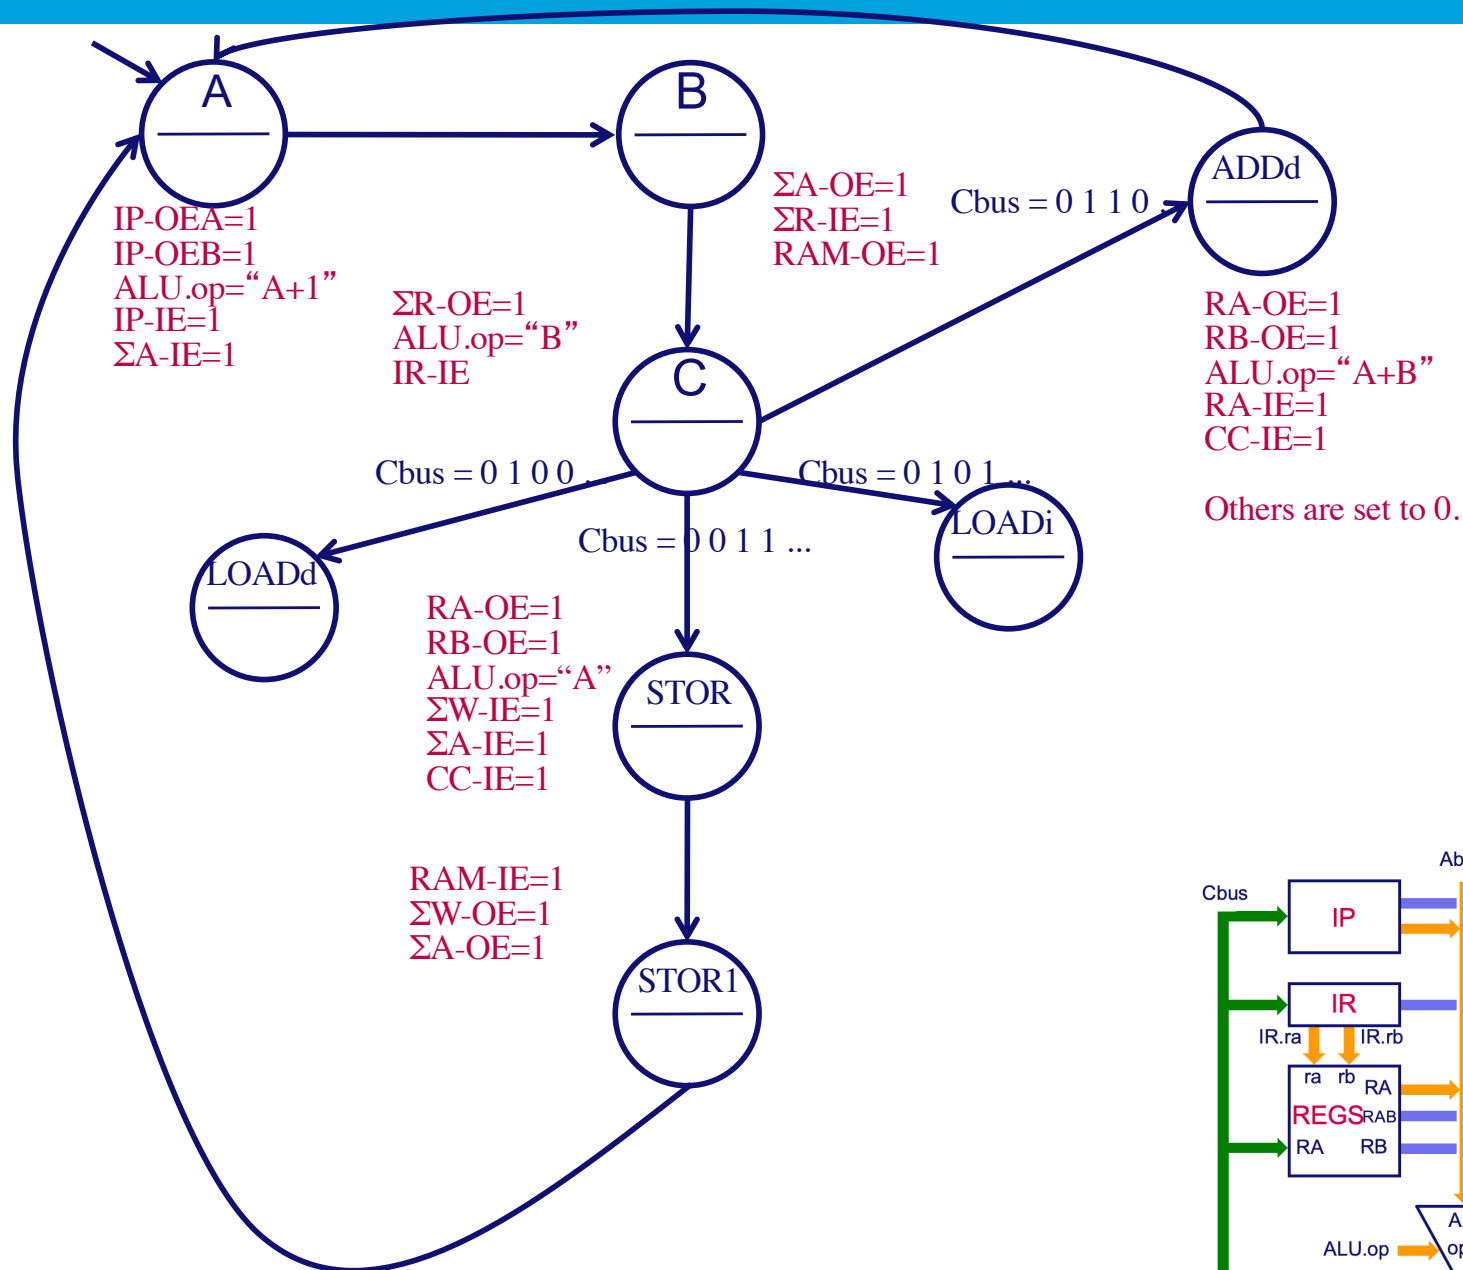
After the last step the Conductor returns to state 0; this is left *implicit*.

Execution of this instruction requires 4 clock cycles, 3 of which for IF!

RA is the register in REGS identified by IR.ra;

RB is the register in REGS identified by IR.rb;

Conductor: execute ADD, direct

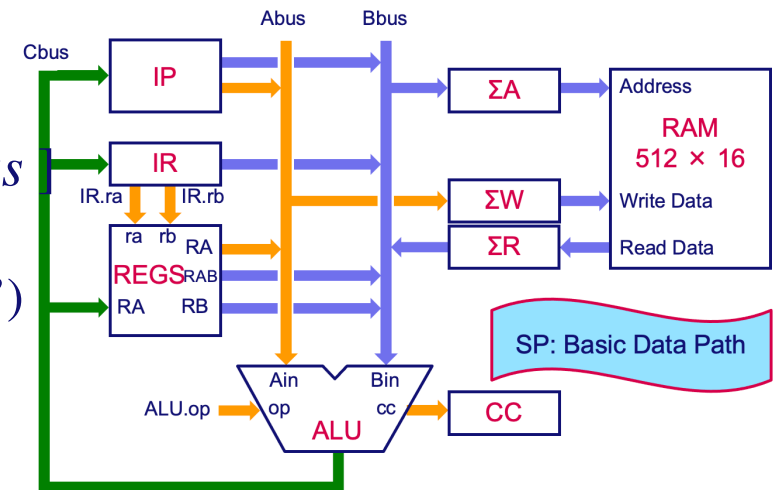


Example: Instruction execution (XOR, indirect)

Instruction: XOR RA [*address*]

Required action: $RA \leftarrow RA \oplus RAM[address]$

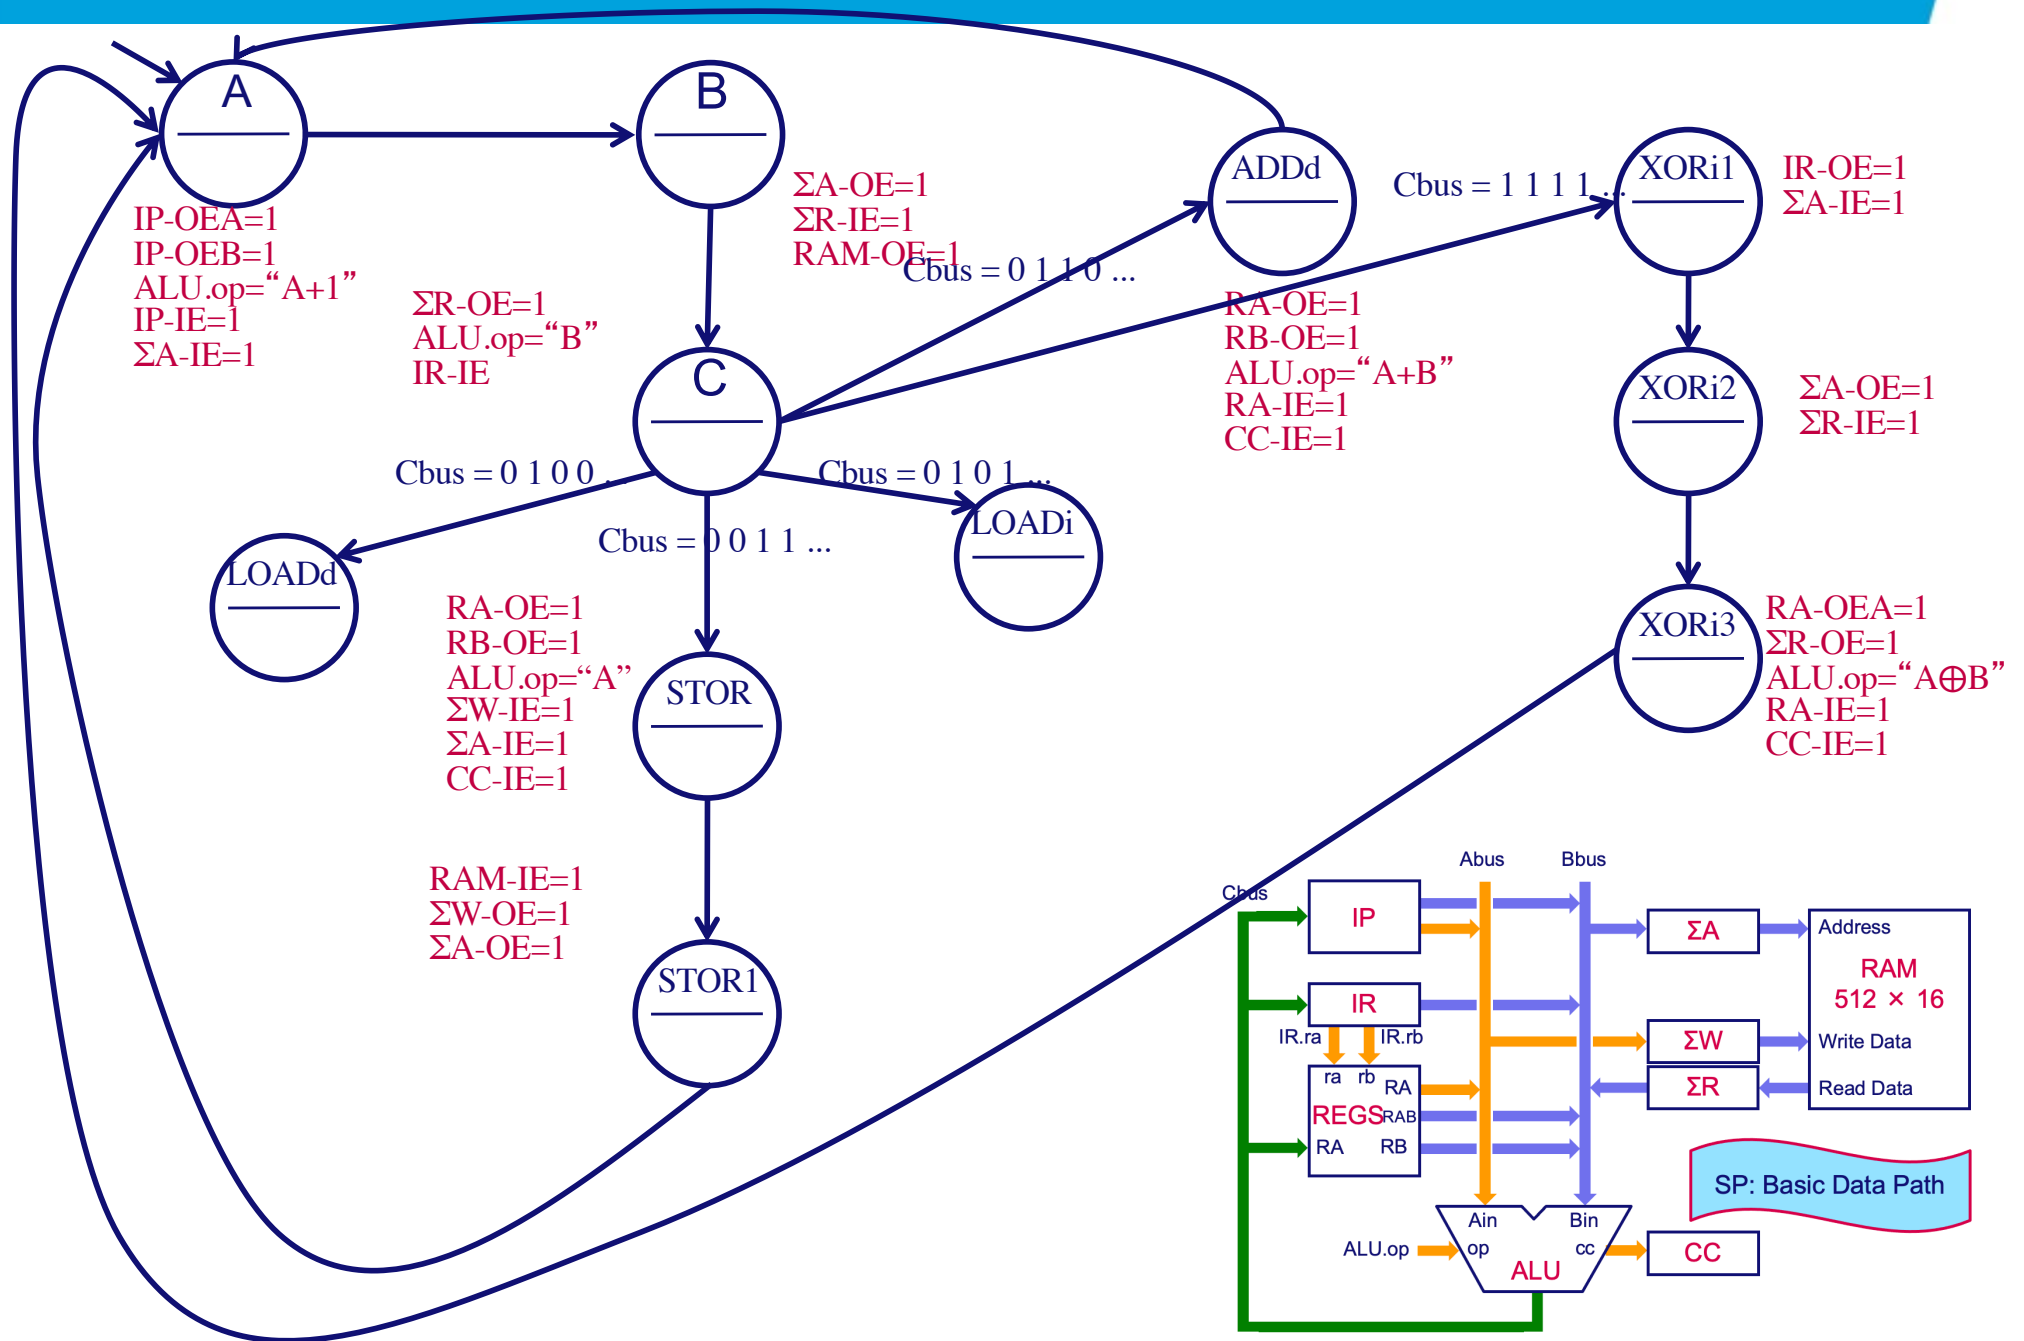
- 0 $\Sigma A, IP \leftarrow IP(Bbus), IP(Abus)+1$ (ALU: “A+1”)
- 1 $\Sigma R \leftarrow RAM[\Sigma A]$
- 2 $IR \leftarrow \Sigma R$ (ALU: “B”)
- 3 $\Sigma A \leftarrow IR.val$ $\Sigma A \leftarrow address$
- 4 $\Sigma R \leftarrow RAM[\Sigma A]$ fetch operand $RAM[address]$
- 5 $RA, CC \leftarrow RA \oplus \Sigma R$ (ALU: “ $A \oplus B$ ”), ALU.cc



Again, after the last step the Conductor returns to state 0: 6 cycles used.

Two additional cycles –compared to “ADD RA RB”– for the *operand fetch*.

Conductor: execute XOR, indirect



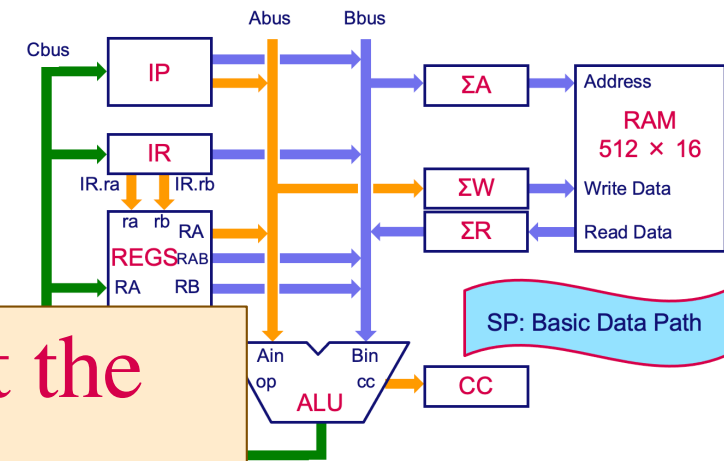
Example: instruction execution (BEQ)

Instruction: BEQ *disp*, case branch taken ($Z=1$)

Required action: $IP \leftarrow IP + disp$

- 0 $\Sigma A, IP \leftarrow IP(Bbus), IP(Abus)+1$ (ALU: “A+1”)
- 1 $\Sigma R \leftarrow RAM[\Sigma A]$
- 2 $IR \leftarrow \Sigma R$ (ALU: “B”)
- 3 $IP \leftarrow IP(Abus) + IR.comp$ (ALU: “A+B”)

Interpret the instruction and the Z flag.



After the last step the Conductor returns to state 0: 4 cycles used.

IR.comp is the sign extended value. This is needed for negative jumps, if the memory is larger than 512 bytes large.

Instruction format: conditional jumps

Instruction format: 0 0 0 1 c c c d d d d d d d d d d



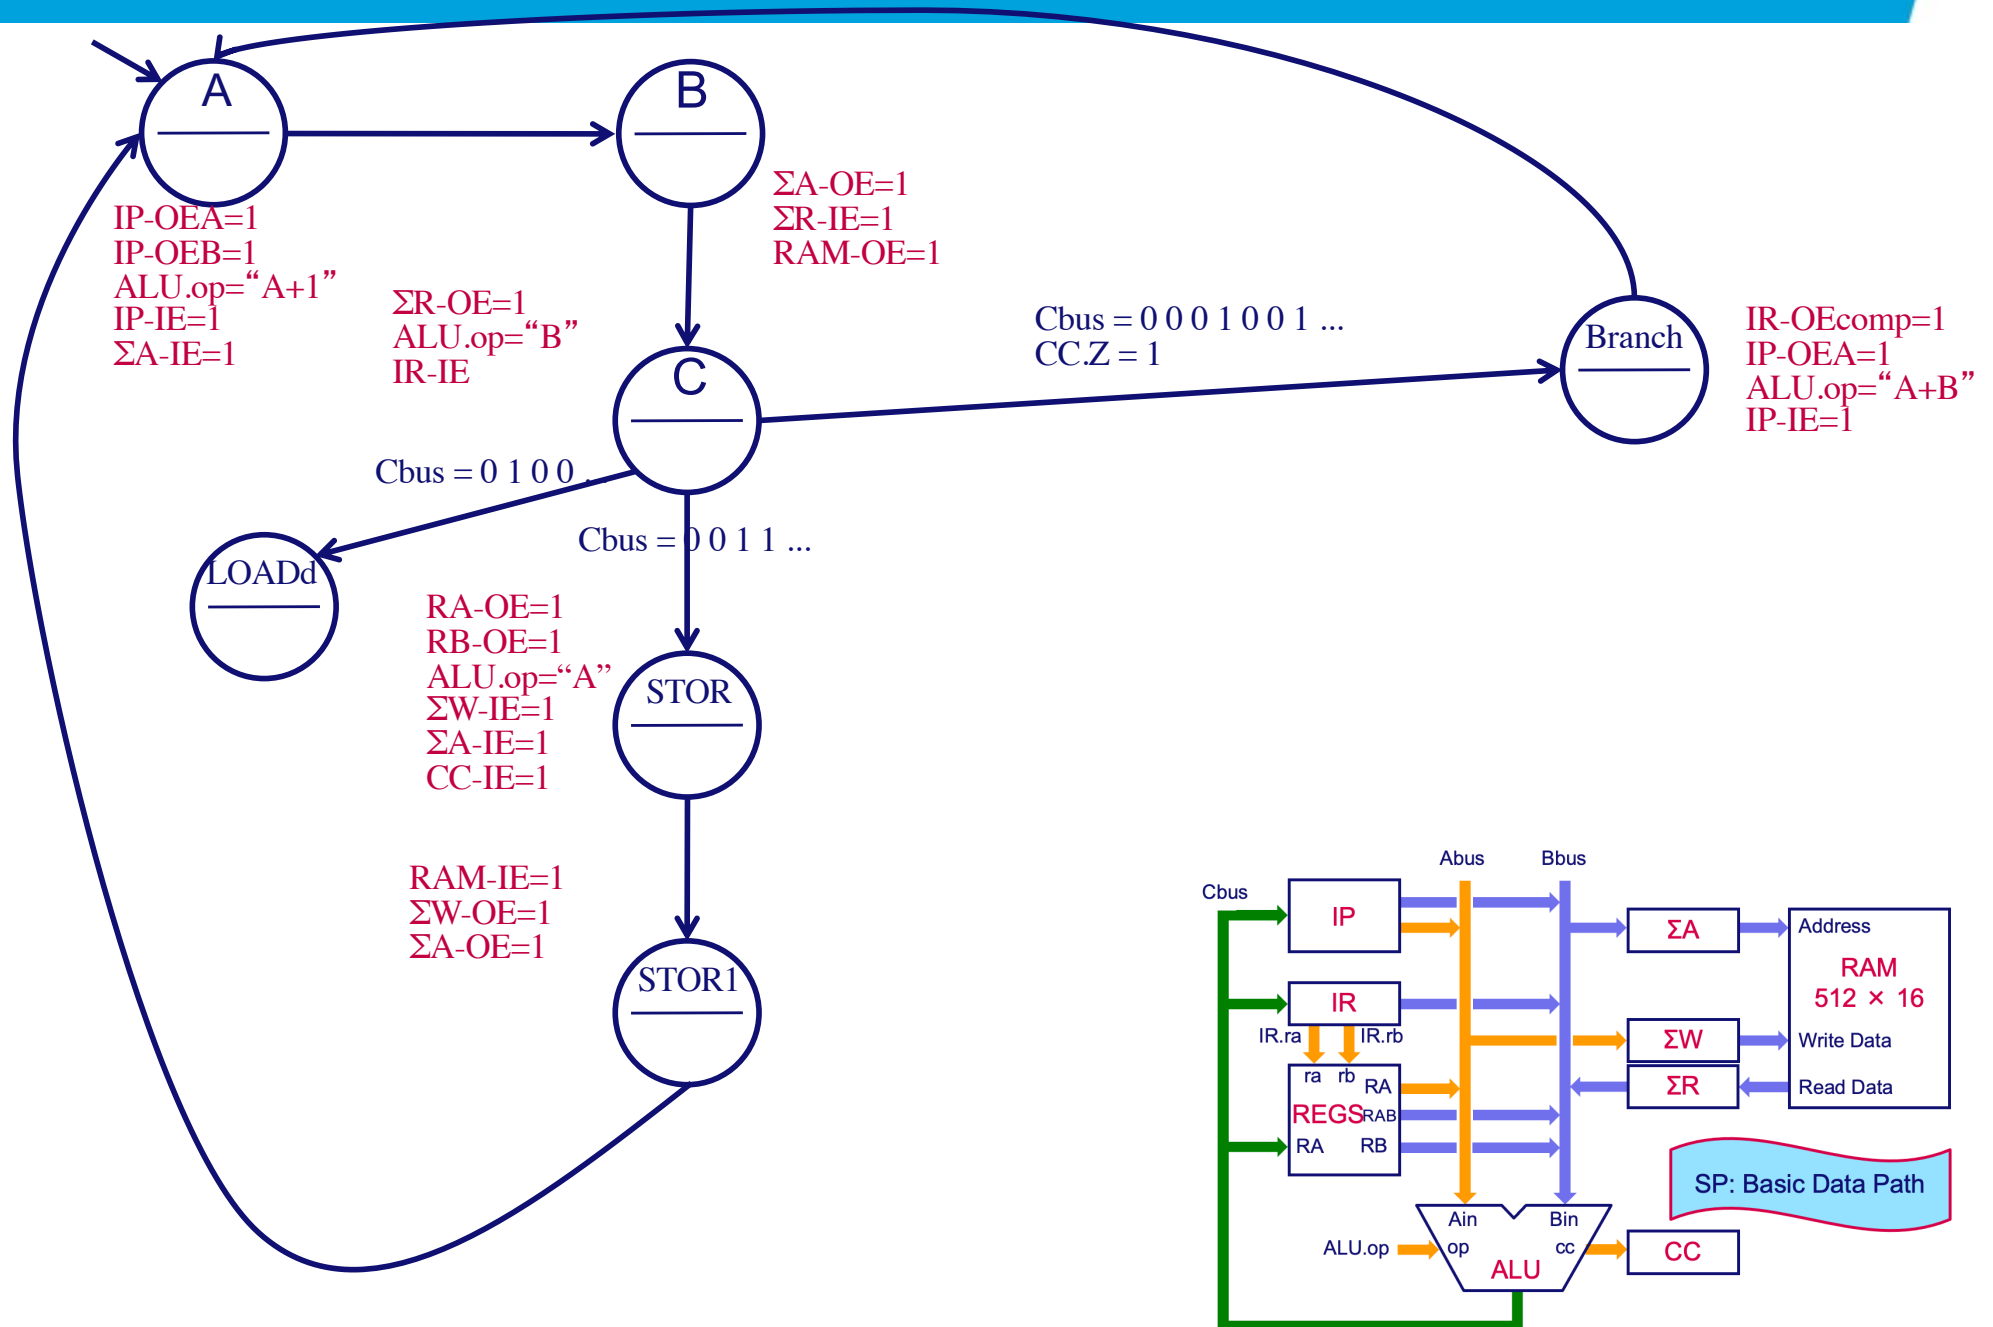
16 bits

BNE	0 0 0	Z = 0
BEQ	0 0 1	Z = 1
BPL	0 1 0	N = 0
BMI	0 1 1	N = 1
BVC	1 0 0	V = 0
BVS	1 0 1	V = 1
BCC	1 1 0	C = 0
BCS	1 1 1	C = 1

BEQ	-37
	0001 001 111011011
BVC	171
	0001 100 010101100

Code relocation requires
relative jumps.

Conductor: execute jump to new location



Example: instruction execution (BEQ)

Instruction: BEQ *disp*, case branch *not* taken ($Z=0$)

Required action: “do nothing, just continue”

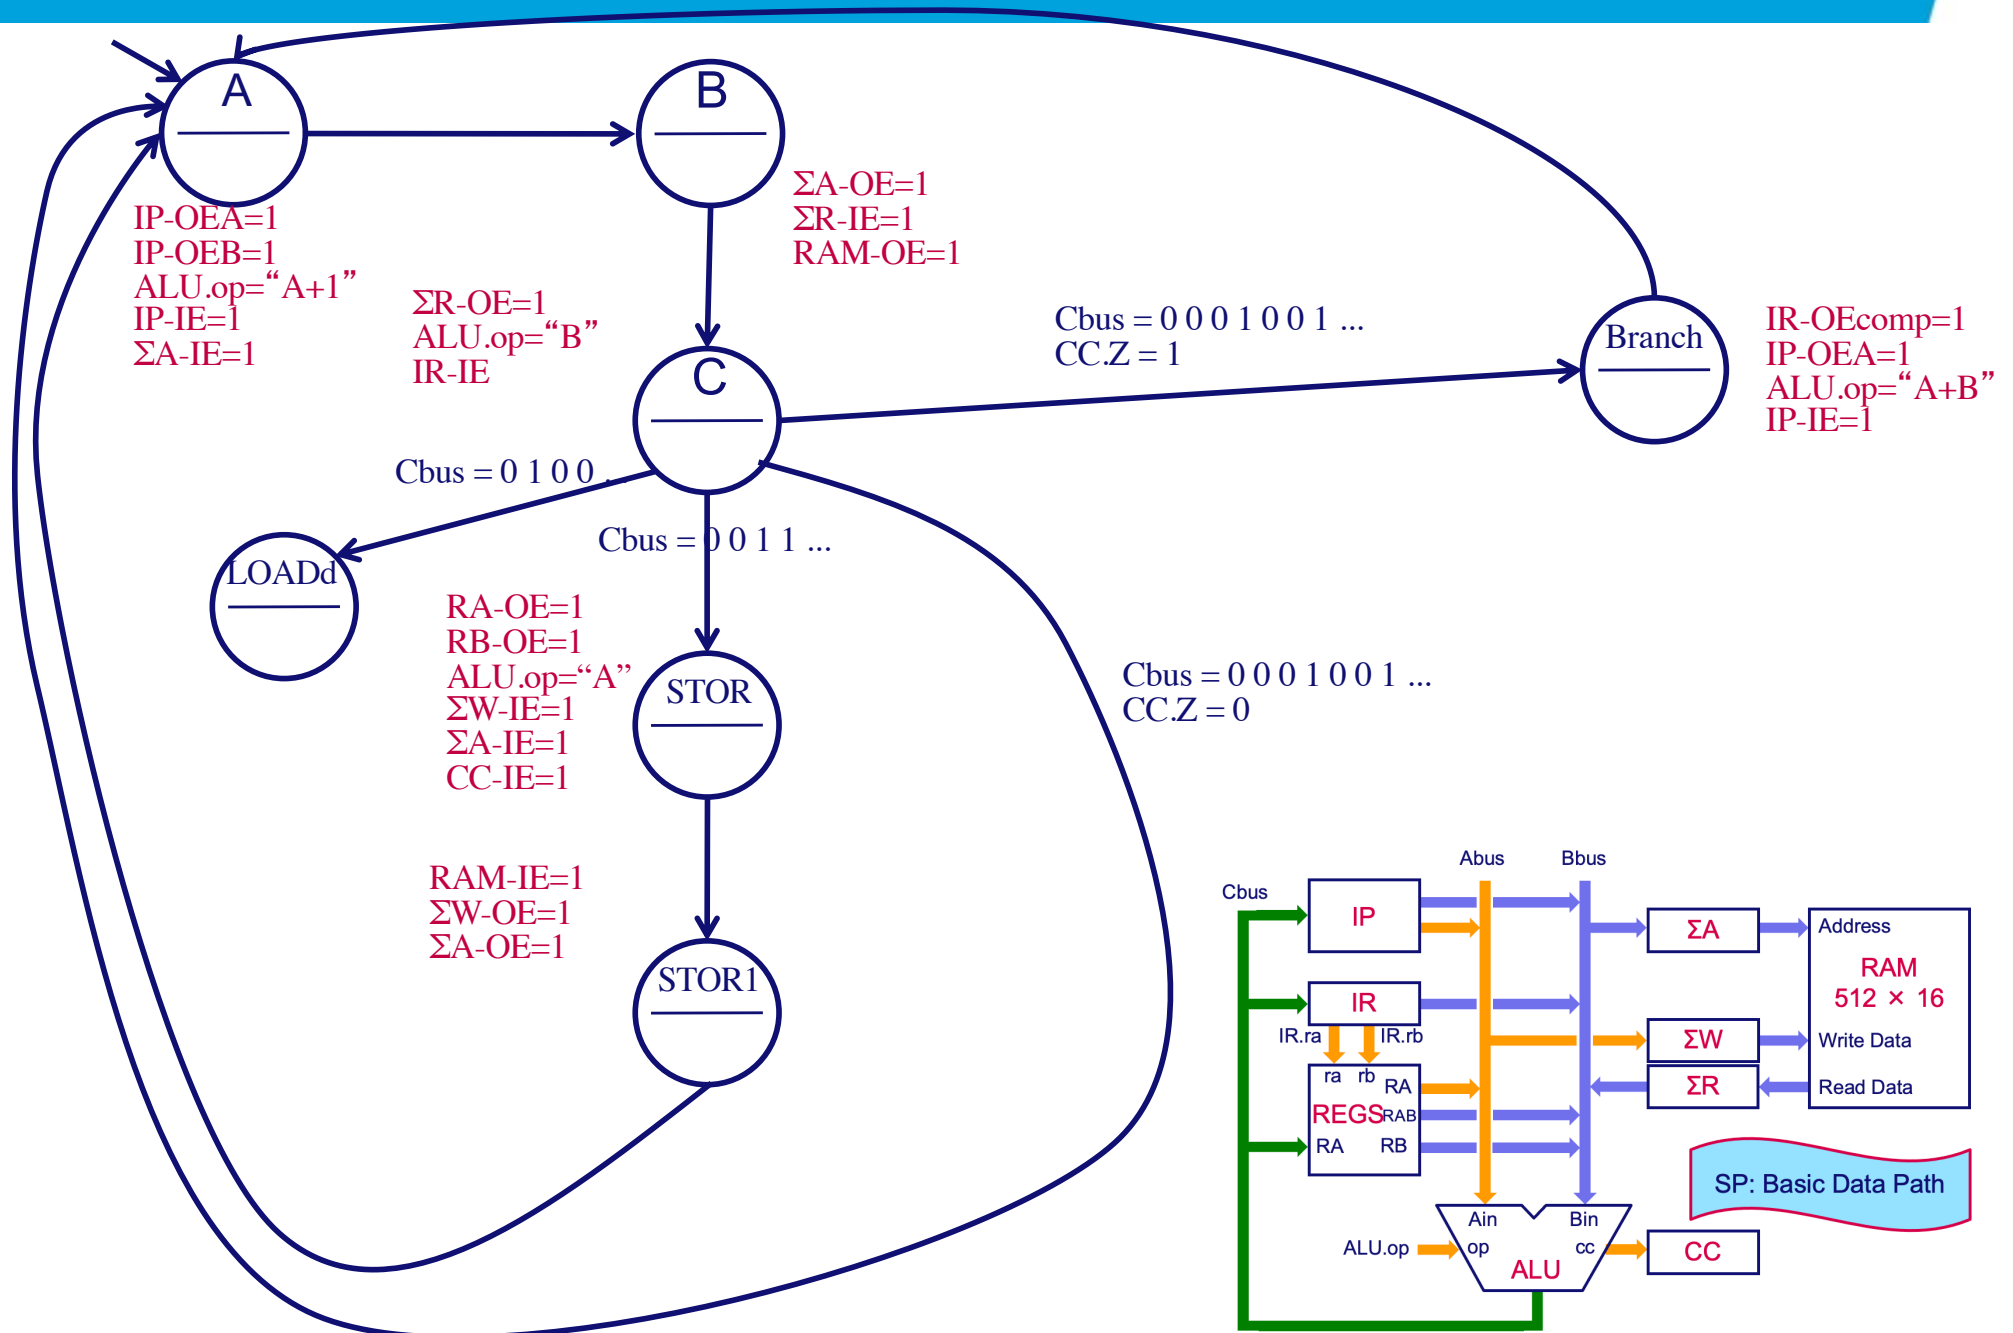
- 0 $\Sigma A, IP \leftarrow IP(Bbus), IP(Abus)+1$ (ALU: “A+1”)
- 1 $\Sigma R \leftarrow RAM[\Sigma A]$
- 2 $IR \leftarrow \Sigma R$ (ALU: “B”)

Interpret the instruction and the Z flag.

After step 2 the Conductor returns to state 0: 3 cycles used.

Hence, in this case “do nothing” really means doing nothing!

Conductor: skip jump to new location




Example: Instruction execution (BRS)

Instruction: BRS *reg displacement*

Required action: $reg, RAM[reg-1], IP \leftarrow reg-1, IP, IP + displacement$


- 0 $\Sigma A, IP \leftarrow IP(Bbus), IP(Abus)+1$ (ALU: “A+1”)
- 1 $\Sigma R \leftarrow RAM[\Sigma A]$
- 2 $IR \leftarrow \Sigma R$ (ALU: “B”)
- 3 $RA \leftarrow RA(Abus)-1$
- 4 $\Sigma A, \Sigma W \leftarrow RAB(Bbus), IP(Abus)$
- 5 $IP, RAM[\Sigma A] \leftarrow IP(Abus)+IR.compl(Bbus), \Sigma W$



Micro programming

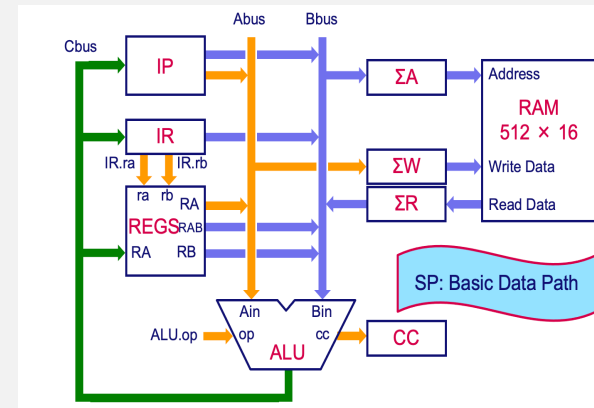

Program in
memory

LOAD R1 1
ADD R1 [R2]
BEQ else_case
....



Microprogram:

- 3 $RA \leftarrow RA (Abus)-1$
- 4 $\Sigma A, \Sigma W \leftarrow \text{RAB} (Bbus), IP (Abu$
- 5 $IP, RAM [\Sigma A] \leftarrow IP (Abus)+IR.$

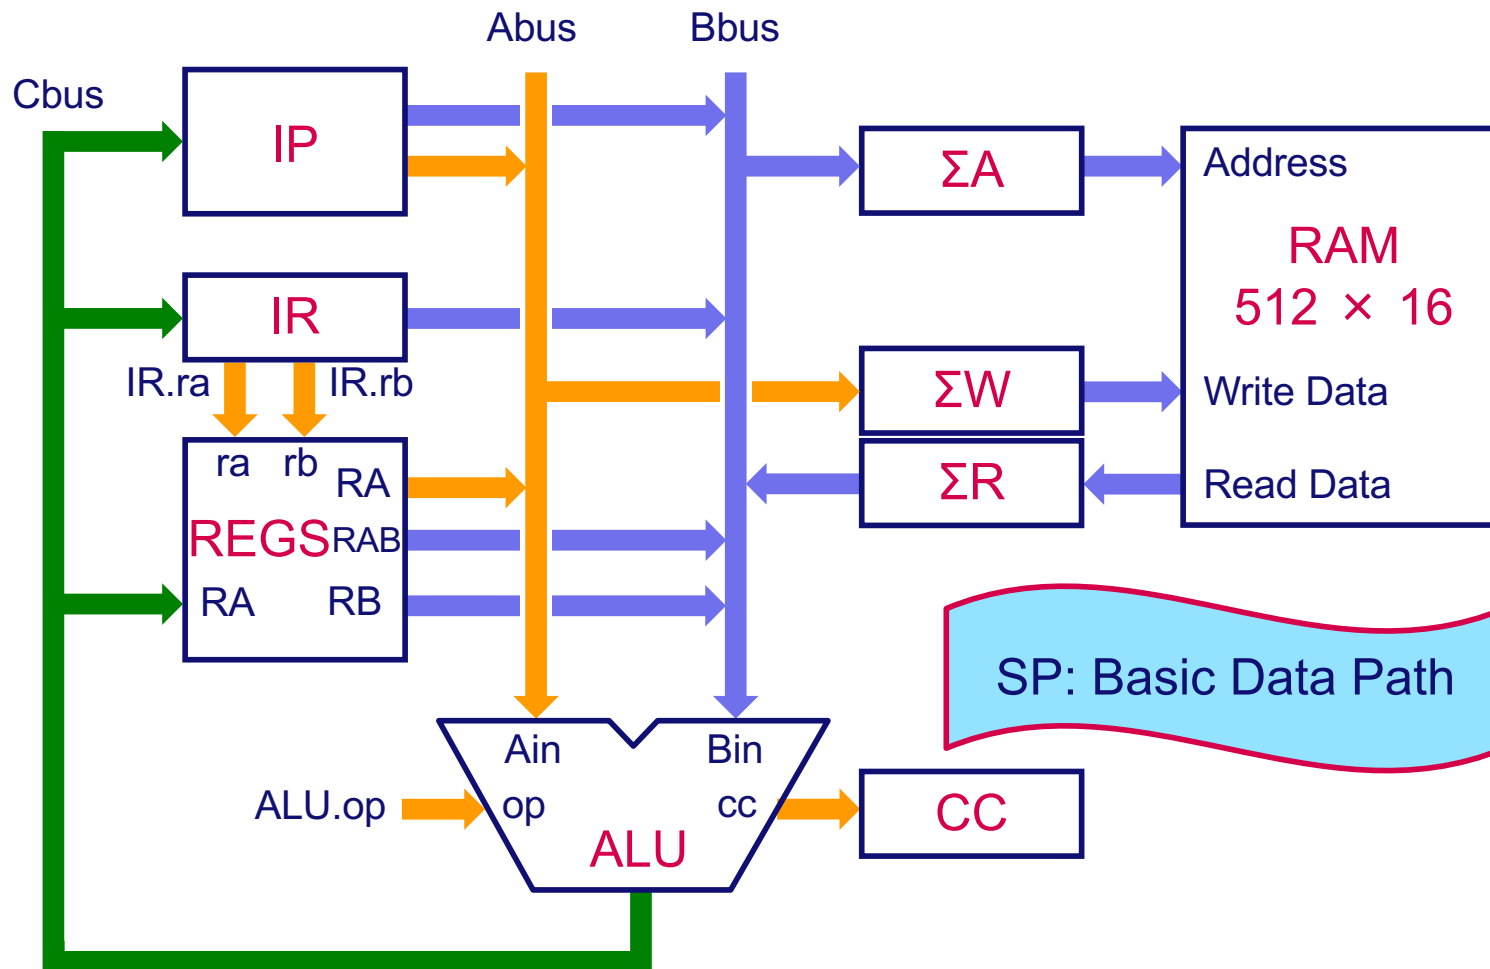


Alternative for
encoding a state
machine.

Questions?

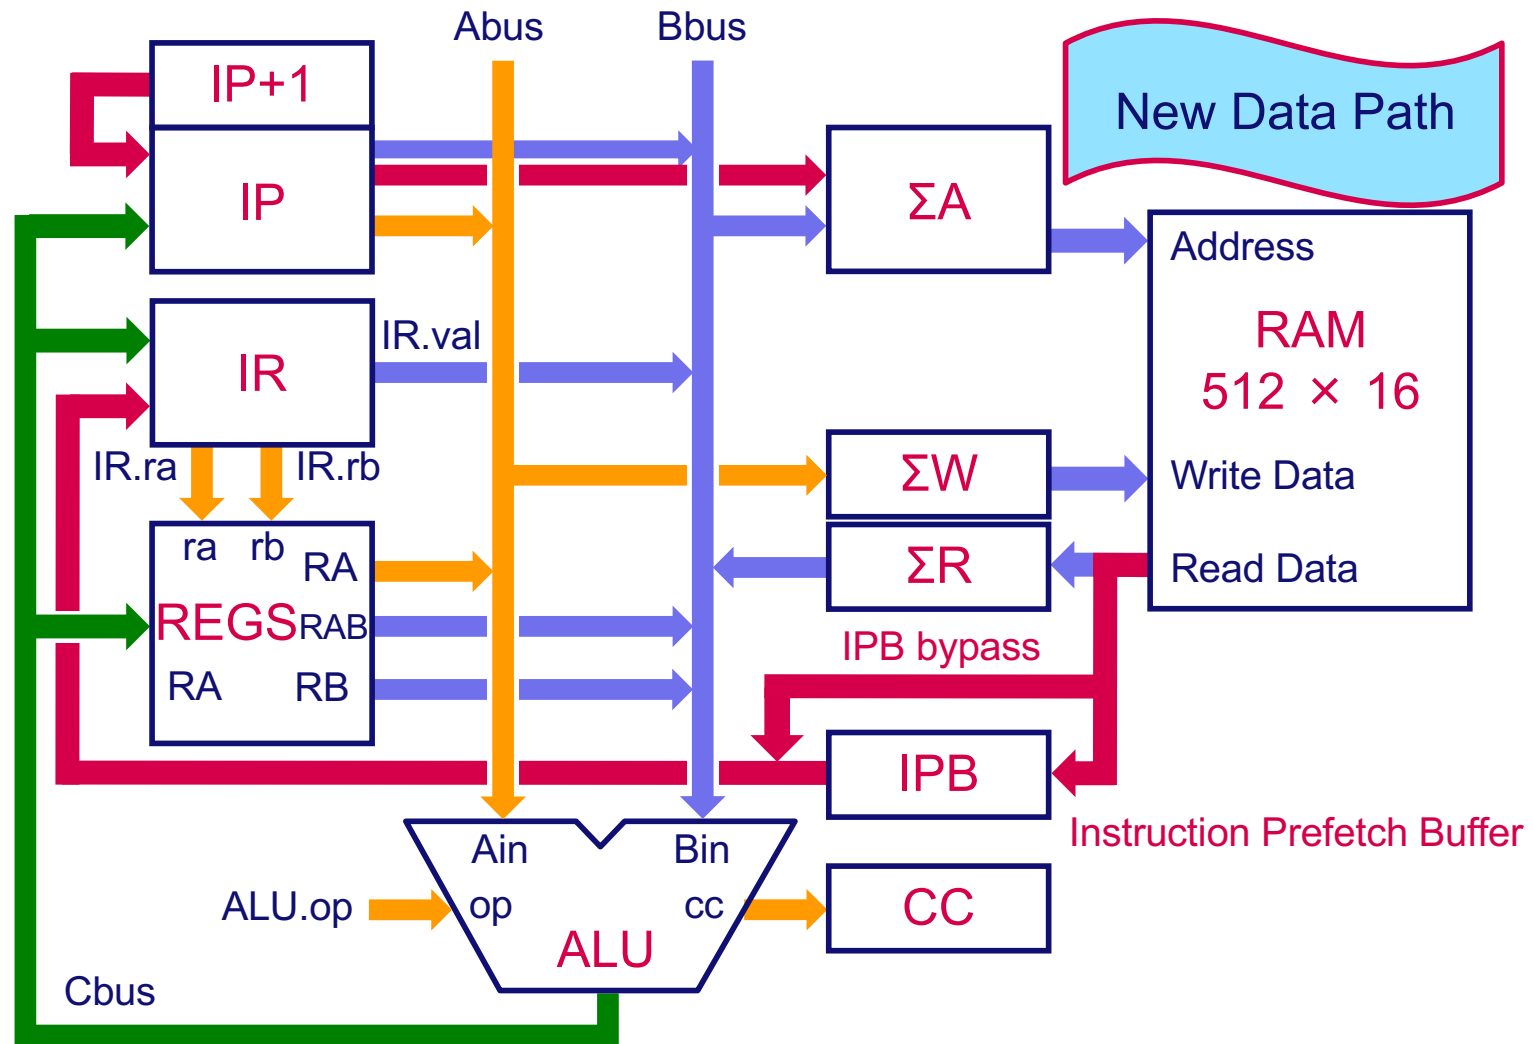


Simple processor: data path



- Abus: IP or RA
- Bbus: IP, IR, ΣR, RB or RAB
- REGS: 2 Reads and 1 Write, simultaneously, $ra[1..0]$ and $rb[1..0]$ select registers

Simple processor: speeding up the machine



- IPB: allows prefetching without overwriting IR too early

Example: faster instruction execution (1)

Instruction: ADD RA RB

Required action: $RA \leftarrow RA + RB$

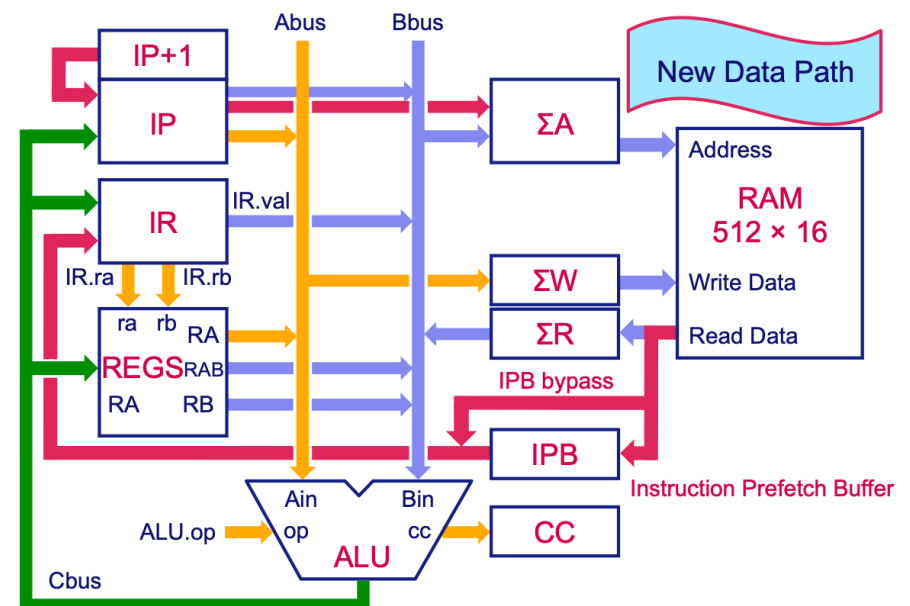
- 0 $\Sigma A, IP \leftarrow IP, IP+1$ using the IP-incrementer
- 1 $IR, \Sigma A, IP \leftarrow RAM[\Sigma A], IP, IP+1$ IR via the IPB-bypass
- 2 $RA, CC, IR, \Sigma A, IP \leftarrow$
 $RA + RB$ (ALU: “A+B”), ALU.cc, $RAM[\Sigma A], IP, IP+1$

In state 1 the value of IP is already sent to ΣA , to initiate prefetching.

In state 2 the next instruction arrives; because the current execution is completed here, it can be placed into IR directly: IPB is not needed here.

In addition, due to the prefetch: IP points 2 instructions *ahead!*

After state 2 the Conductor goes to a new instruction, with a prefetched instruction. Thus, instructions of this type need only 1 cycle, instead of 4 !



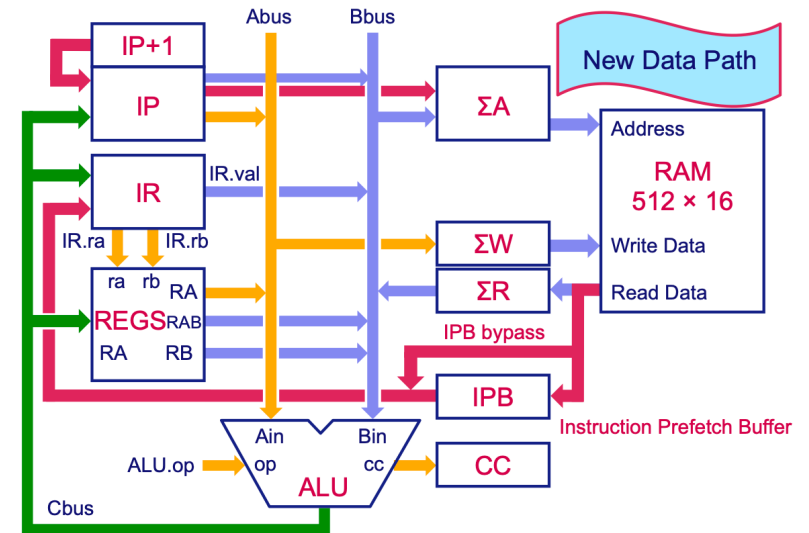
Example: faster instruction execution (2)

Instruction: XOR RA [*address*]

Required action: $RA \leftarrow RA \oplus RAM[address]$

- 0 $\Sigma A, IP \leftarrow IP, IP+1$
- 1 $IR, \Sigma A, IP \leftarrow RAM[\Sigma A], IP, IP+1$ fetch instruction, start prefetch
- 2 $\Sigma A, IPB \leftarrow IR.val, RAM[\Sigma A]$ start operand fetch, save pref. instr. in IPB
- 3 $\Sigma R \leftarrow RAM[\Sigma A]$ fetch operand
- 4 $RA, CC, IR, \Sigma A, IP \leftarrow$ execute, prefetch from IPB
 $RA \oplus \Sigma R$ (ALU: " $A \oplus B$ ") , ALU.cc , IPB, IP , IP+1

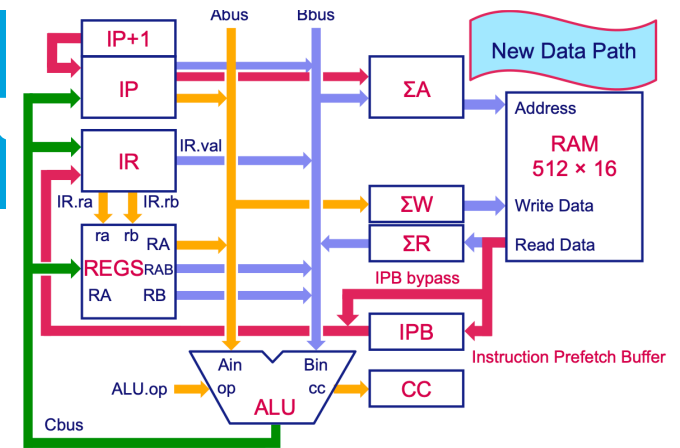
With a prefetch instruction starts in 2,
ends in 4: 3 cycles used.



Example: faster instruction execution

Instruction: **STOR RA [RB]**

Required action: $\text{RAM} [\text{RB}] \leftarrow \text{RA}$



- 0 $\Sigma A, IP \leftarrow IP, IP+1$
- 1 $IR, \Sigma A, IP \leftarrow \text{RAM} [\Sigma A], IP, IP+1$ fetch instruction
- 2 $\Sigma A, \Sigma W, CC, IPB \leftarrow RB, RA, \text{ALU.cc (ALU: "A")}, \text{RAM} [\Sigma A]$
- 3 $\text{RAM} [\Sigma A], IR, \Sigma A, IP \leftarrow \Sigma W, IPB, IP, IP+1$ fetch from IPB

Tricky: In step 3 the Memory is busy with “ $\text{RAM} [\text{RB}] \leftarrow \text{RA}$ ”; hence, the instruction prefetch is initiated in 2; IPB is used as a buffer to delay the next instruction fetch.

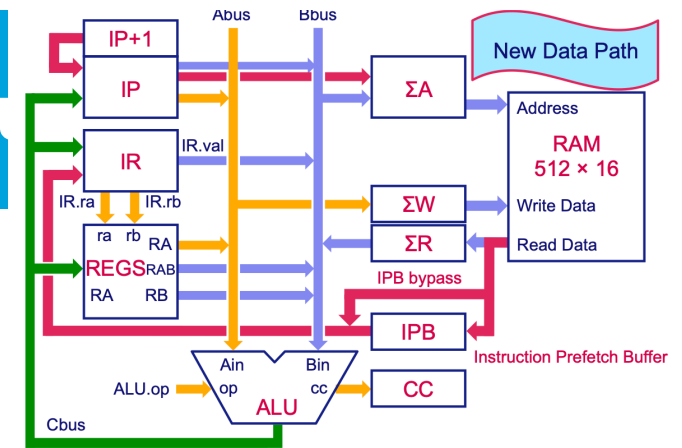
With a prefetch the processor starts at 2: it only needs two instructions.

Example: faster instruction execution

Instruction: BEQ *disp*, case branch taken ($Z=1$)

Required action: $IP \leftarrow IP + disp$

- 0 $\Sigma A, IP \leftarrow IP, IP+1$
- 1 $IR, \Sigma A, IP \leftarrow RAM[\Sigma A], IP, IP+1$ initiate prefetch: will be wasted!
 $\{ IR = \text{"BEQ } disp", IP's \text{ value is } 1 \text{ too large}, Z = 1 \}$
- 2 $IP \leftarrow IP(Abus) - 1$ (ALU: "A-1") correct IP's value
- 3 $IP \leftarrow IP(Abus) + IR.comp$ (ALU: "A+B") perform the jump
- 4 $\Sigma A, IP \leftarrow IP, IP+1$ initiate *target* IF
- 5 $IR, \Sigma A, IP \leftarrow RAM[\Sigma A], IP, IP+1$ fetch *target* instruction



Prefetching the *next* instruction now is part of executing the *current* one.

Hence, again the Conductor goes back to state 2: 4 cycles for this case.

Warning: this is a very tricky game!

Alternative: to *prevent* IP correction an additional *buffer register* may be used, or the ALU may be extended with an *additional operation*: "A+B-1".

Example: Faster instruction execution (4b)

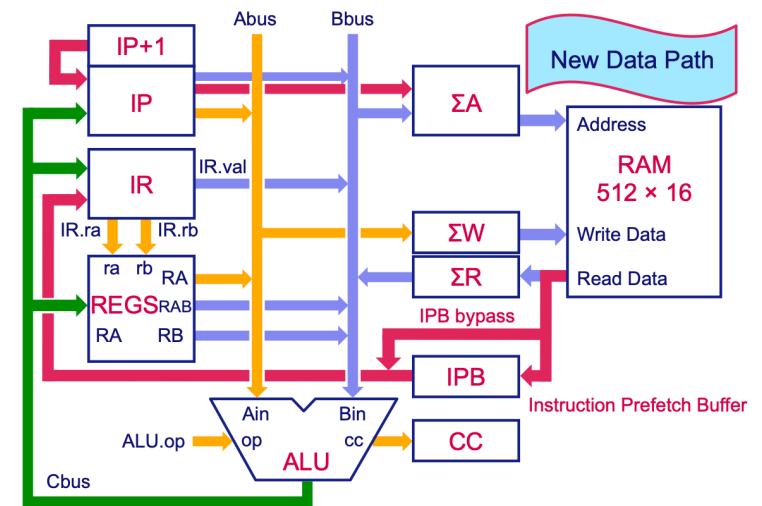
Instruction: BEQ *disp* , case branch *not* taken ($Z=0$)

Required action: “do nothing, just continue”

- 0 $\Sigma A, IP \leftarrow IP, IP+1$
- 1 $IR, \Sigma A, IP \leftarrow RAM[\Sigma A], IP, IP+1$ initiate prefetch: not wasted!
{ $IR = \text{“BEQ } disp\text{”}, Z = 0$ }
- 2 $IR, \Sigma A, IP \leftarrow RAM[\Sigma A], IP, IP+1$ fetch *next* instruction

Because fetching the next instruction has been initiated already in step 1,
“do nothing” entails no more than completing the Instruction Fetch in step 2.

With a prefetch, the executing the instruction starts
in state 2: only 1 cycle now.



Simple Processor: Summary

Instruction Execution Times, without and with prefetching

Instruction:	no prefetch:	prefetch:	memory accesses:
ADD R2 R3	4	1	1
XOR R2 [<i>address</i>]	6	3	2
STOR R0 [R3]	5	2	2
BEQ <i>disp</i>	4 / 3	4 / 1	1

A substantial improvement!

Other possibilities for speeding up:

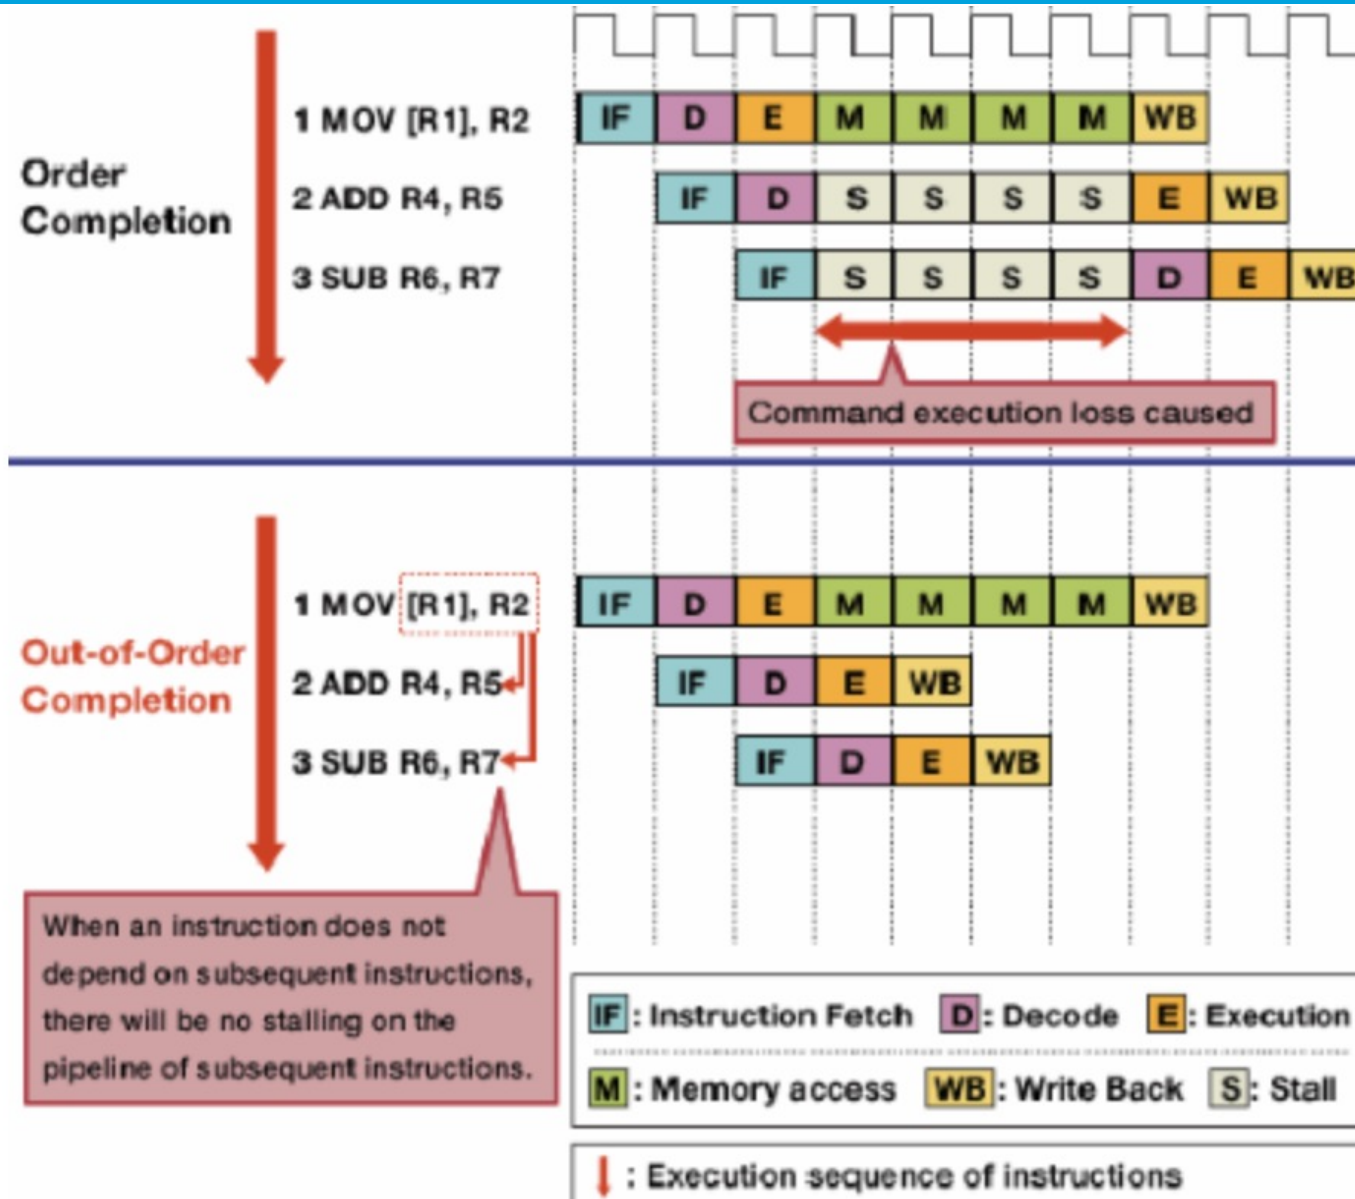
- adding more buffer registers; example: IP, to retain “old” value
- adding more bypasses; example: from ALU to ΣA .
- adding dedicated adders; example: address calculations.

The number of memory accesses needed is a *lower bound!*

Out of order execution.

On parallel computers changing the ordering can be harmful.

Hence, **barrier** instructions that explicitly prevent out of order execution.

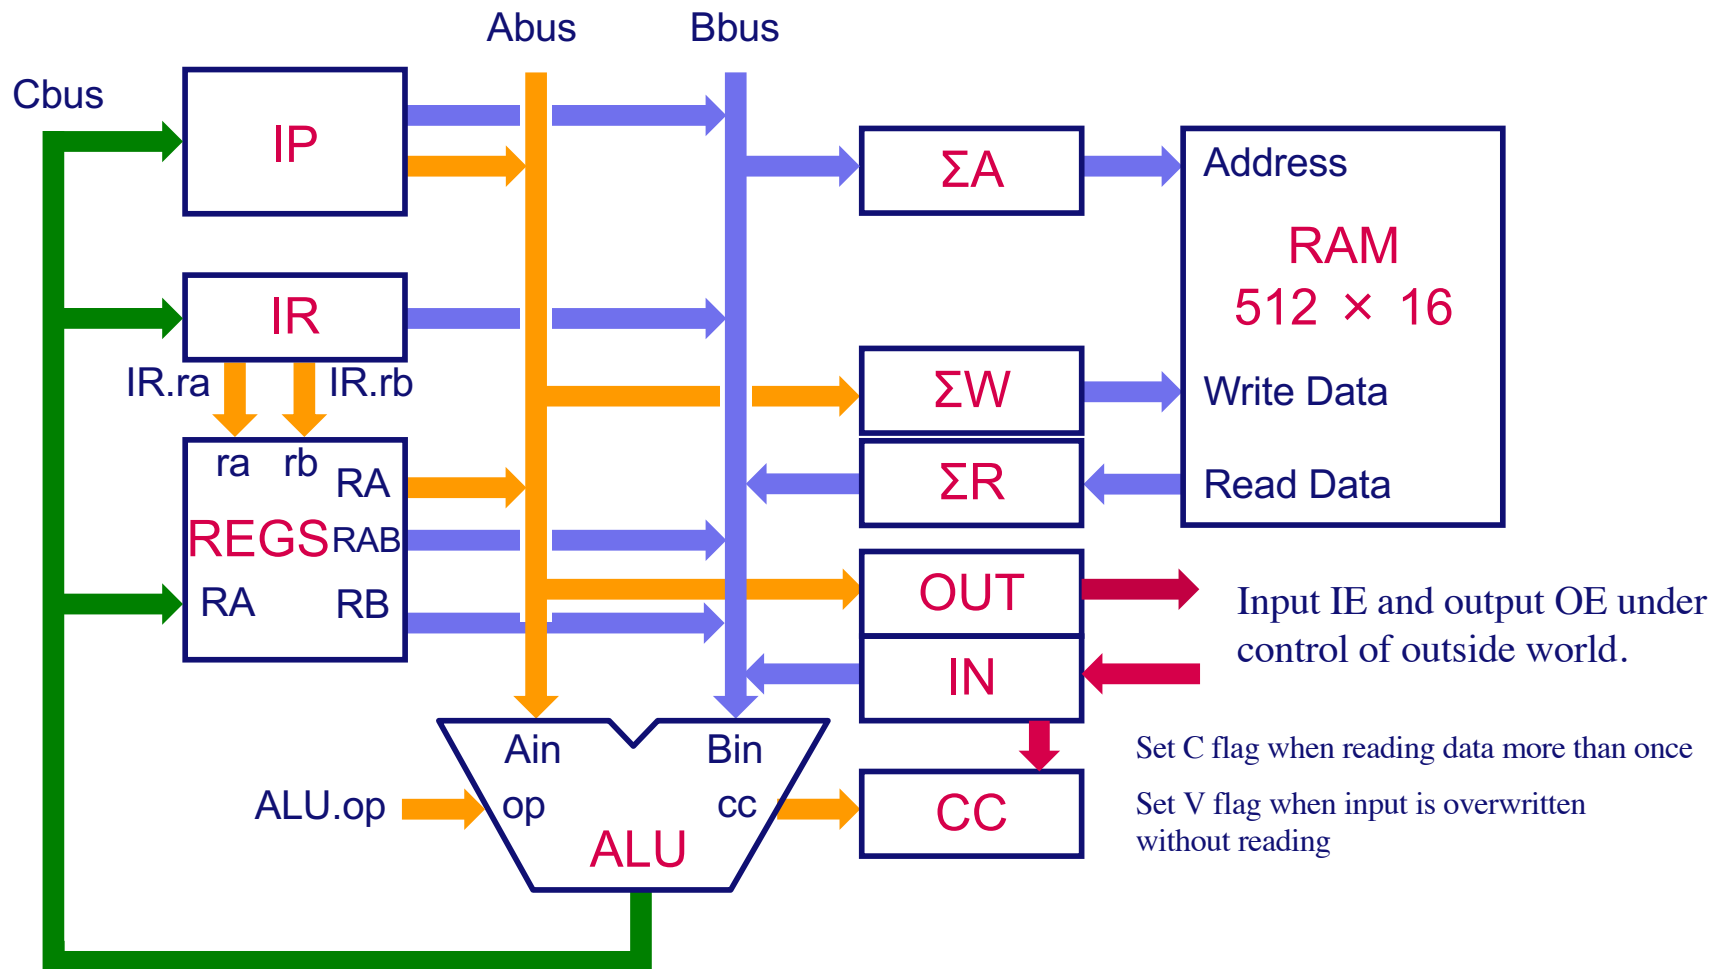


Picture from: renesasrulz.com

Questions?



Simple processor: Input and output.



IN register 0 0 1 0 r r 1 0 0 0 0 0 0 0 0 0

OUT register 0 0 1 0 r r 0 1 0 0 0 0 0 0 0 0

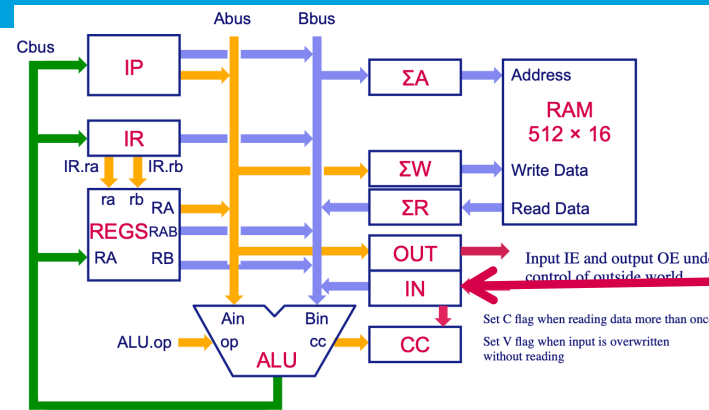
Input and output: Polling.

Count keyboard inputs in register R1.

```

LOAD R1 0;      Reset key stroke count.
start_keyboard_read:
  IN R0;         Read input buffer.
  BVS overwrite_error;
                  Data entry in buffer has
                  been missed.
  BCS start_keyboard_read;
                  If carry flag is set, there is
                  no new data, so loop back.

  ADD R1 1;      Count the key stroke.
  BRA start_keyboard_read;
overwrite_error:
                  This is the place for error handling.
    
```



QWERTY KEYBOARD

~	!	@	#	\$	%	^	&	*	()	+	=	Delete
Tab	Q	W	E	R	T	Y	U	I	O	P	{	}	\
Caps	A	S	D	F	G	H	J	K	L	:	"	'	Enter
Shift	Z	X	C	V	B	N	M	<	>	?/			Shift
Ctrl		Alt									Alt		Ctrl

http://www.computerhope.com

Polling
computationally
expensive.

Summary

What did we learn:

- We implemented all instructions on the simple data path.
- It is possible to change the ‘simple data path’, add buses and registers if that would be useful.
- We can optimize the execution time of the instructions, for instance by prefetching instructions.