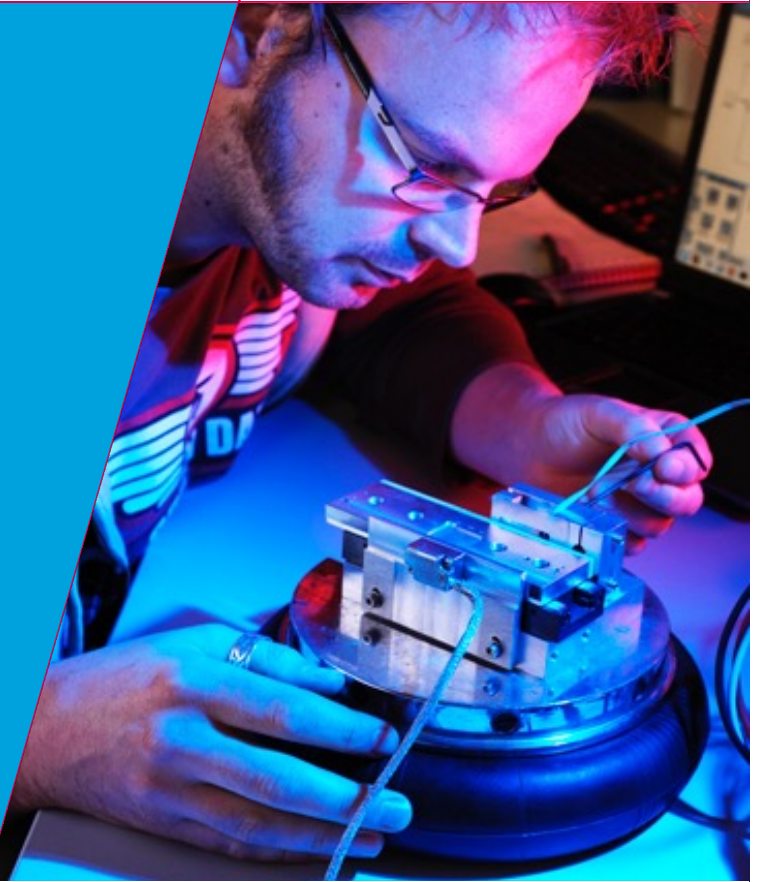


2IC30: Computer systems Translating Higher Languages

Jan Friso Groote



TU / **e**

Technische Universiteit
Eindhoven
University of Technology

Where innovation starts

A program in machine code (PP2 processor)

Another Example

```
;
;   INPUT:  R0 = N      , precondition: 0 <= N
;   OUTPUT: R1 = Fib(N) , R2 = Fib(N+1)
;   USES:   R3
;
```

Machine code

```
001000100 000000000
001001000 000000001
010100000 000000000
000000000 000000101
001001100 100000001
001000100 100000010
001101000 100000011
000011101 000010100
010000000 000000001
000000011 111111010
000100010 111110001
```



; Rob Hoogerwoord.

;

; INPUT: R0 = N , precondition: $0 \leq N$

; OUTPUT: R1 = Fib(N) , R2 = Fib(N+1)

; USES: R3

;

;

LOAD R1 0 ; invariant: $m = N - R0$ and
LOAD R2 1 ; $R1 = \text{Fib}(m)$ and $R2 = \text{Fib}(m+1)$
CMP R0 0 ; test R0, to prepare the condition codes
BRA +5 ; enter the "loop"
LOAD R3 R1 ; $R2 = \text{Fib}(m+1)$ and $R3 = \text{Fib}(m)$
LOAD R1 R2 ; $R1 = \text{Fib}(m+1)$ and $R2 = \text{Fib}(m+1)$ and $R3 = \text{Fib}(m)$
ADD R2 R3 ; $R1 = \text{Fib}(m+1)$ and $R2 = \text{Fib}(m+2)$
TRCS ; abort when unsigned overflow, i.e., carry set
SUB R0 1 ; $R0 := R0 - 1$, i.e.: $m := m + 1$, hence:
; $R1 = \text{Fib}(m)$ and $R2 = \text{Fib}(m+1)$
BNE -6 ; if $R0 = 0$ then $R1 = \text{Fib}(N)$
RTS



@END

In PP2 assembly

Assembly Language (1)

Problem

- Manually counting branch distances still is tedious and error prone
- Manually calculating memory addresses is tedious and error prone

Solution

- Use names to identify the targets of branch instructions: *labels*

Assembly Language (1)

In ARM assembly

```
; Rob Hoogerwoord;
;   INPUT:  R0 = N      , precondition: 0 <= N
;   OUTPUT: R1 = Fib(N) , R2 = Fib(N+1)
;   USES:   R3
;
;
Fibonacci : MOV    R1,#0      ; invariant: m = N - R0 and
                MOV    R2,#1      ; R1 = Fib(m) and R2 = Fib(m+1)
                CMP    R0,#0      ; test R0, to prepare the condition codes
                B      Fibon_while ; enter the "loop"
Fibon_do   : MOV    R3,R1      ; R2 = Fib(m+1) and R3 = Fib(m)
                MOV    R1,R2      ; R1 = Fib(m+1) and R2 = Fib(m+1) and R3 = Fib(m)
                ADDS   R2,R2,R3    ; R1 = Fib(m+1) and R2 = Fib(m+2)
                SWIOV                ; software interrupt when (unsigned) overflow
                SUBS   R0,#1      ; R0 := R0 - 1, i.e.: m := m + 1 , hence:
                                ; R1 = Fib(m) and R2 = Fib(m+1)
Fibon_while : BNE    Fibon_do    ; if R0 = 0 then R1 = Fib(N)
Fibonacci_end : MOV PC,LR      ; return to invoking program
```

Memory Addressing (0)

- Suppose an integer variable is located at word 13

- Loading this integer into register R1 (say):

```
MOV R0, #13
```

```
LDR R1, [ R0 ]
```

- This is still error prone! Therefore: *always* use names!
- Calling the variable *x* (say), it can be *bound* to its location in the data segment by means of an “EQU” pseudo-instruction:

```
.equ x 13 ; now x equals 13
```

```
.set x 13 ; equivalent to .equ x 13
```

```
x = 13 ; also equivalent to .equ x 13
```

- The same instruction to load *x* into R1:

```
MOV R0, #x
```

```
LDR R1, [ R0 ]
```

- Advantages: less error prone and better readable

Does not work for arbitrary *x*.
Why?

Memory Addressing (0)

- Calling the variable x (say), it can be *bound* to its location in the data segment by means of an “EQU” pseudo-instruction:

.equ x 134657463; now the address of x equals 134657463

.set x 134657463

$x = 134657463$

- The same instruction to load x into R1 (leads to an error by the assembler):

MOV R0, # x

LDR R1, [R0]

- The same instruction to load x into R1:

LDR R0, = x

LDR R1, [R0]

Example of a concrete assembly program.

```
.global main
```

```
.data
```

```
.equ constant 45
```

```
var_x:      .int      10
```

```
var_y:      .byte     'A', 0x31, 32, 0x33, 34, $35
```

```
string:     .asciz    "This is a tekst.\n"
```

```
main:      LDR R0, =var_x
```

```
           LDR R3, [R0]
```

```
           STR R3, [SP, #-1]!
```

```
           BL  subroutine
```

```
           STR R4, [SP], #1
```

```
           STR R3, [R4, #constant]
```

```
           B   main
```

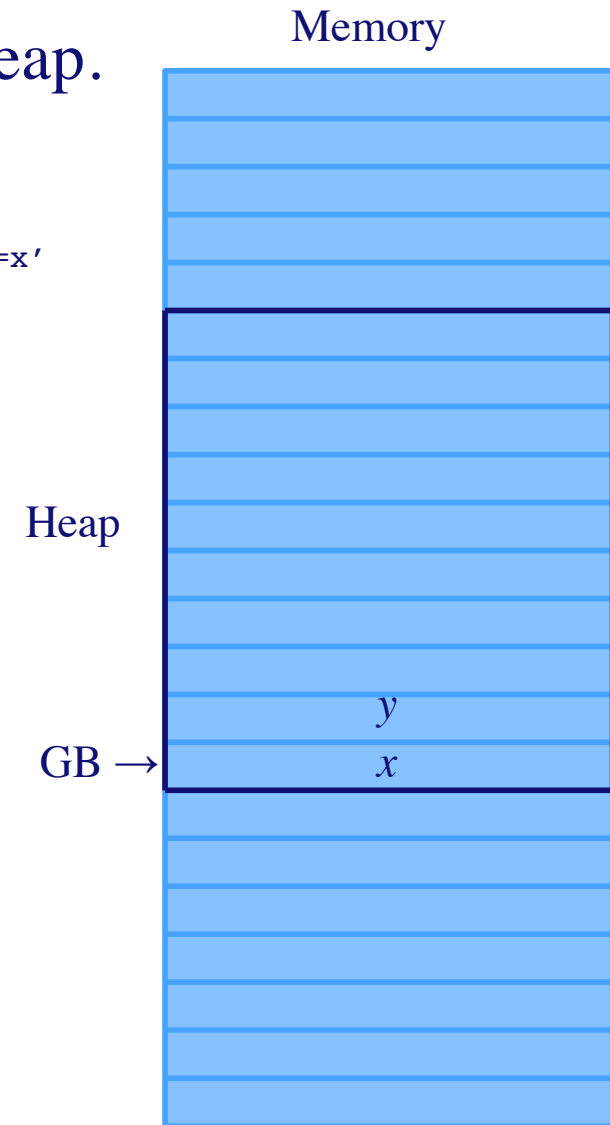
```
.end
```


Example of the use of the heap (0)

Swap values of variables x and y on the heap.

```
; INPUT: x, y
; OUTPUT: y is old value of x, x is old value of y: x=y', y=x'
;
global main:
main:
    LDR    GB, =global_base
    LDR    R0, [GB,#x]        ; R0:=x
    LDR    R1, [GB,#y]        ; R1:=y
    STR    R0, [GB,#y]        ; y:=R0
    STR    R1, [GB,#x]        ; x:=R1

.data
.equ x 0          ; x resides at GB+0
.equ y 1          ; y resides at GB+1
.equ global_base .... ; definition of the global base
```



Example of the use of the heap (1)

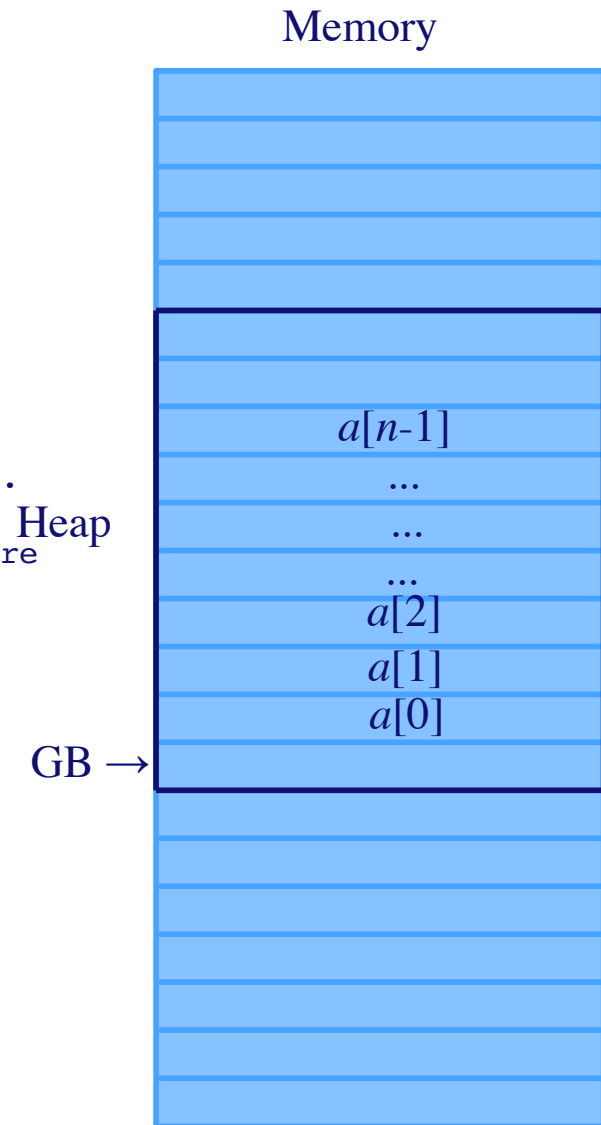
Reset an array $a[0,...,n-1]$.

```
; INPUT: R0, a; Precondition R0=n
; OUTPUT: Postcondition a[i]=0 for all 0<=i<n.
;
.global main

main:
reset_a :      MOV    R1, #0           ; R1:=0 R1 represents i.
               LDR     R3, =global_base
               ADD     R3, #a           ; R3:=GB+a, address where
                                       ; array a starts.
while_reset_a: MOV    R2,R1
               SUBS    R2,R0           ; R2:=(R1-R0) R2=i-n
               BEQ     endwhile_reset ; if i=n, stop
               MOV     R4,#0           ; R4:=0
               STORE   R4 [R3,R1]      ; a[i]:=0
               ADD     R1,#1           ; i:=i+1
               B       while_reset_a   ; loop back

endwhile_reset: ...

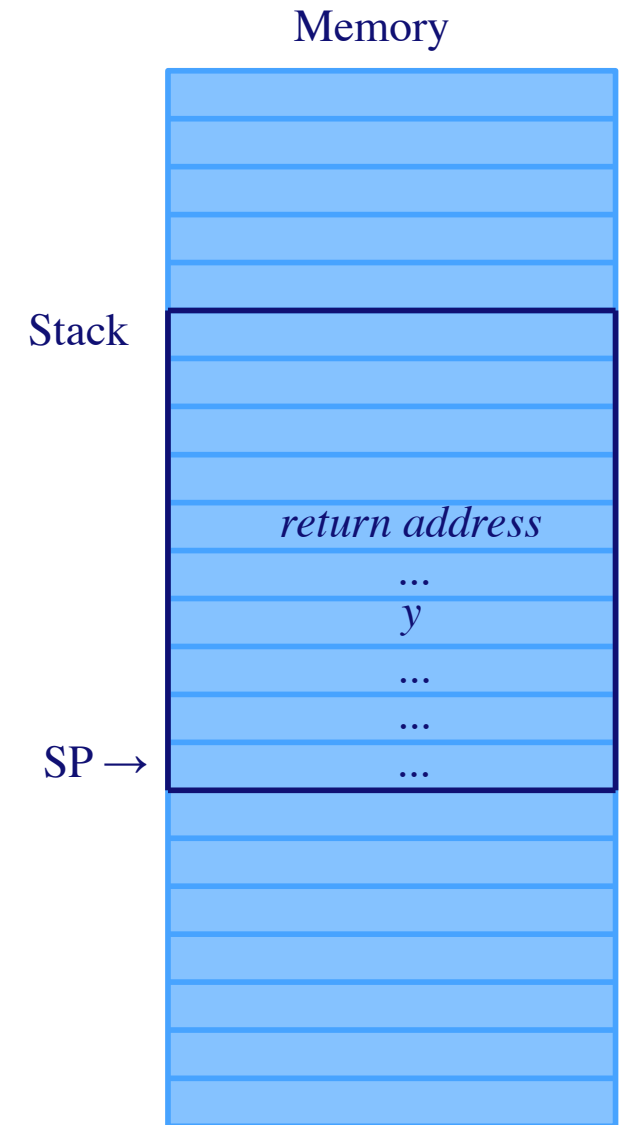
.data
.equ a ...           ; base address of a
.equ global_base ... ; definition of global base
```



Memory Addressing (1)

Local variables: used only temporarily

- Local variables can be kept in registers, or ... on the *stack*!
- The stack grows *downwards* in memory; to *allocate* -- i.e. reserve -- 5 words on the stack, use: `SUB SP 5`.
- The addresses of these 5 words are `SP`, `SP+1`, ..., `SP+4`.
- Better is to name them; if we use the word at position 3 as a local variable `y` (say), it can be *bound* to its location in the Stack data segment by means of an “.equ” pseudo-instruction: `.equ y 3`; now the address of `y` equals `SP + y`.



Stack addressing: allocating local variables

- ❑ **Example:** A subroutine needs 6 words to store local variables.

procedure *f()*

var *a, b, c, d, e, x*: integer;

.....

.....

end;

- ❑ **Allocation:** By *decreasing* SP by the amount –6– needed.
- ❑ **Deallocation:** By *increasing* SP by the amount –6– used.
- ❑ Allocation takes place at the *beginning* of the subroutine, *before* the *actual body* of the subroutine.
- ❑ Deallocation takes place at the *end* of the subroutine, *after* the *actual body* of the subroutine, but *before* the *return*.
- ❑ Every subroutine must *leave* the stack effectively *unchanged!!!*

Stack addressing: allocating local variables

```
entry:    SUB    SP    6    ; allocate 6 words for locals
          .equ   a    0    ; "declare" local variable a
          .equ   b    1    ; "declare" local variable b
          ...
          .equ   x    5    ; "declare" local variable x
          ...
          "actual body of the subroutine"
          ...
exit:     ADD     SP    6    ; release the 6 words...
          RTS      ; ... and return
```

Generation of programming languages

First generation: Machine code

Second generation: Assembly code

Third generation: Higher level languages, such as ALGOL 60, Fortran, Cobol, ALGOL 68, PASCAL, C++, Java.

Prof.dr. Frans Kruseman Aretz:

“With the third level languages we thought our troubles were over. We could not be more wrong...”



Foto: TU/e

Generation of programming languages

First generation: Machine code

Second generation: Assembly code

Third generation: Higher level languages, such as ALGOL 60, Fortran, Cobol, ALGOL 68, PASCAL, C++, Java.

Fourth generation: Not so clear. Domain specific languages???
SQL, HTML?

FORTH

APL

Large(r) Projects: Program Design

- *Separate* program design and construction of the Assembly Language program: *firstly*, design and write the program in Java, C++, ... ; *secondly*, translate that program into Assembly Language.
- Why???

Large(r) Projects: Program Design

- *Separate* program design and construction of the Assembly Language program: *firstly*, design and write the program in Java, Pascal, ..., or in a DSL (Domain Specific Language) or just plain mathematics; *secondly*, translate that program into Assembly Language.
- Why???
- Humans are very bad in handling complexity. Always work as abstract as as possible. Use mathematics, or any domain specific notation....

Example with recursion:

$n!$ with a recursive procedure, $n > 0$:

```
procedure fac(var n:integer);  
var h:integer;  
begin  
  if (n≠1) then begin  
    h:=n;  
    n:=n-1;  
    fac(n);  
    n:=h*n;  
  end  
end;
```

In assembler:

CMP: subtract without storing the results.
MULS: multiply arguments.

```
procedure fac(var n:integer);  
var h:integer;  
begin  
  if (n≠1) then begin  
    h:=n;  
    n:=n-1;  
    fac(n);  
    n:=h*n;  
  end  
end;
```

```
fac_main: CMP R0 1  
          BEQ fac_exit  
          STOR R0 [--SP]  
          SUB R0 1  
          BRS fac_main  
fac_ret:  LOAD R1 [SP++]  
          MULS R0 R1  
fac_exit: RTS
```

Note: parameter n is passed in register R0.
Local variable h is stored on the stack.

Questions?

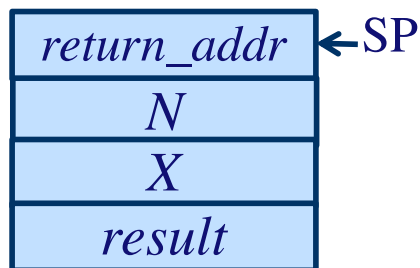


Example: Translating a Java function into Assembly (1)

```
int PowerSum(int X, int N)  
  // requires:  $0 \leq N$   
  // result:  $\text{PowerSum}(X, N) = (\sum i : 0 \leq i \leq N : X^i)$   
  { int x, y, z, n ;  
    z = 0 ; y = 1 ; x = X ; n = N ;  
    while (n != 0)  
    { while (n % 2 == 0)  
      { y = y * (1 + x) ; x = x * x ; n = n / 2 ; }  
      z = z + y ; y = y * x ; n = n - 1 ;  
    }  
    return z ;  
  }
```

Example: Translating a Java function into Assembly (2)

- (0) Decide how the *parameters* will be transferred: parameters X and N are *values* to be transferred *to* the subroutine, and the function's *result* is to be transferred back *from* the subroutine to the calling program:
- These three (integer) values are passed via the *Stack*:
 - Just *before* the call of the subroutine, the calling program:
 - reserves 1 word on the stack for the function's *result*;
 - then, it pushes X onto the stack;
 - next, it pushes N onto the stack;
 - finally, it calls the subroutine.
 - Directly *after* the call of the subroutine, the calling program:
 - retrieves 1 the function's *result* from the reserved word for this purpose;
 - then, it *cleans up* the stack;



```
int PowerSum(int X, int N)
// requires:  $0 \leq N$ 
// result:  $PowerSum(X, N) = (\sum i : 0 \leq i \leq N : X^i)$ 
```

Example: Translating a Java function into Assembly (3)

; Calling sequence to invoke subroutine PowerSum:

```

call_begin : SUB    SP    1          ; allocate one word on the stack,
                                           ; for the result.

              LOAD   R0    X          ; here we just write "X": may
                                           ; be more complicated.

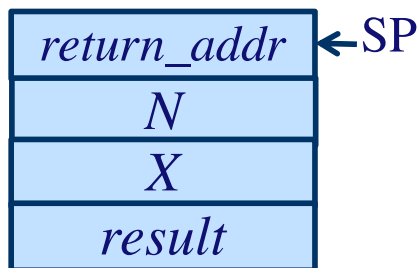
              STOR   R0    [--SP]     ; push X onto the stack.
              LOAD   R0    N          ; here we just write "N".
              STOR   R0    [--SP]     ; push N onto the stack.
              BRS    PowerSum         ; call subroutine PowerSum.

call_end :   LOAD   R0    [SP+2]      ; retrieve PowerSum's result...
              STOR   R0    z          ; ... and store it somewhere,
                                           ; here just called "z"

              ADD    SP    3          ; clean up Stack: deallocate the 3
                                           ; words used

              ...                    ; and continue...

```



Example: Translating a Java function into Assembly (4)

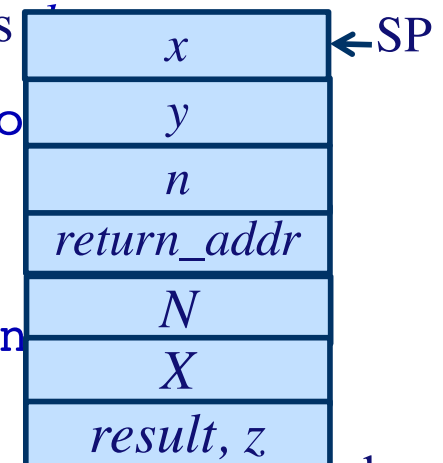
(1) When execution of the subroutine starts, the top of the stack has this

RAM[SP] = "return address of the calling pro

RAM[SP+1] = N

RAM[SP+2] = X

RAM[SP+3] = "word for the (anonymous) function



(2) Now, decide where *local variables* x, y, z, n will be allocated. Because the final value of z will be returned as the result, z may be identified with the result word at SP+3. Variables x, y and n are put on the stack:

$x = \text{RAM}[\text{SP}+0]$

$y = \text{RAM}[\text{SP}+1]$

$n = \text{RAM}[\text{SP}+2]$

$z = \text{RAM}[\text{SP}+3+3]$

Compensate for two local variables on the stack.

int PowerSum(int X, int N)

// requires: $0 \leq N$

// result: $\text{PowerSum}(X, N) = (\sum i : 0 \leq i \leq N$

{ int x, y, z, n ;

$z = 0$; $y = 1$; $x = X$; $n = N$;

while ($n \neq 0$)

Example: Translating a Java function into Assembly (1)

```
int PowerSum(int X, int N)  
// requires:  $0 \leq N$   
// result:  $\text{PowerSum}(X, N) = (\sum i : 0 \leq i \leq N : X^i)$   
{ int x, y, z, n ;  
  z = 0 ; y = 1 ; x = X ; n = N ;  
  while (n != 0)  
  { while (n % 2 == 0)  
    { y = y * (1 + x) ; x = x * x ; n = n / 2 ; }  
    z = z + y ; y = y * x ; n = n - 1 ;  
  }  
  return z ;  
}
```

Example: Translating a Java function into Assembly (5)

<pre> .equ local_vars 3 .equ PowerSum_N 1+1 .equ PowerSum_X 2+1 .equ PowerSum_z 3+1 .equ Powersum_x 0 .equ PowerSum_y 1 .equ PowerSum_n 2 PowerSum_begin: </pre>	<pre> z = 0 ; y = 1 ; x = X ; n = N ; while (n != 0) { while (n % 2 == 0) { y = y * (1 + x) ; x = x * x ; n = n / 2 ; } z = z + y ; y = y * x ; n = n - 1 ; } </pre>
---	--

<pre> x = RAM[SP+0] y = RAM[SP+1] n = RAM[SP+2] z = RAM[SP+6] </pre>	<pre> SUB SP local_vars ; reserve room on the stack LOAD R0 0 ; initialize variables STOR R0 [SP+PowerSum_z] ; z ← 0 LOAD R0 1 ; R0 ← 1 STOR R0 [SP+PowerSum_y] ; y ← 1 LOAD R0 [SP+PowerSum_X] ; R0 ← X STOR R0 [SP+PowerSum_x] ; x ← X LOAD R0 [SP+PowerSum_N] ; R0 ← N STOR R0 [SP+PowerSum_n] ; n ← N PowerSum_while0: BEQ PowerSum_od0 ; exit outer repetition ; flags reflect value of n </pre>
--	---

; continued on next page...

Example: Translating a Java function into Assembly (6)

```
PowerSum_while1: LOAD R0 [SP+PowerSum_n] ; R0 ← n
                  AND  R0  1                ; R0 ← n%2
                  BNE   PowerSum_od1        ; IF not n%2=0, THEN exit
                                          ; inner repetition
```

```
PowerSum_dol:    LOAD R0 [SP+PowerSum_x] ; R0 ← x
                  ADD  R0  1                ; R0 ← x + 1
                  LOAD R1 [SP+PowerSum_y] ; R1 ← y
                  MULS R1  R0                ; R1 ← y × (x+1)
                  STOR R1 [SP+PowerSum_y] ; y ← R1
                  LOAD R0 [SP+PowerSum_x] ; R0 ← x
```

$x = \text{RAM}[\text{SP}+0]$
 $y = \text{RAM}[\text{SP}+1]$
 $n = \text{RAM}[\text{SP}+2]$
 $z = \text{RAM}[\text{SP}+6]$

```

LOAD R0 [SP+PowerSum_n] ; R0 ← n
MULS R0 1                ; R0 ← n%2
STOR R0 [SP+PowerSum_n] ; n ← R0
LOAD R0 [SP+PowerSum_x] ; R0 ← x
SHL R0 1                 ; R0 ← x+1
MULS R1 R0                ; R1 ← y × (x+1)
STOR R1 [SP+PowerSum_y] ; y ← R1
LOAD R0 [SP+PowerSum_x] ; R0 ← x
BR   PowerSum_while1     ; inner repetition

```

$z = 0 ; y = 1 ; x = X ; n = N ;$
 $\text{while } (n \neq 0)$
 $\{ \text{while } (n \% 2 == 0)$
 $\{ y = y * (1 + x) ; x = x * x ; n = n / 2 ; \}$
 $z = z + y ; y = y * x ; n = n - 1 ;$

$\text{; inner repetition}$
 PowerSum_od1:

; continued

Example: Translating a Java function into Assembly (6)

PowerSum_while1:	LOAD	$z = 0 ; y = 1 ; x = X ; n = N ;$
	AND	
	BNE	$\text{while } (n \neq 0)$
		$\{ \text{while } (n \% 2 == 0)$
PowerSum_dol:	LOAD	$\{ y = y * (1 + x) ; x = x * x ; n = n / 2 ; \}$
	ADD	
	LOAD	$z = z + y ; y = y * x ; n = n - 1 ;$
	MULS R1, R0	$; R1 \leftarrow y * (x + 1)$
	STOR R1 [SP+PowerSum_y]	$; y \leftarrow R1$
	LOAD R0 [SP+PowerSum_x]	$; R0 \leftarrow x$
	MULS R0, R0	$; R0 \leftarrow x \times x$
	STOR R0 [SP+PowerSum_x]	$; x \leftarrow R0$
	LOAD R0 [SP+PowerSum_n]	$; R0 \leftarrow n$
	SHIFTRIGHT R0	$; n \leftarrow n / 2$
	STOR R0 [SP+PowerSum_n]	$; n \leftarrow R0$
	BRA PowerSum_while1	$; \text{jump back to beginning of}$
		$; \text{inner repetitionPowerSum_od1}$

x	$=$	$\text{RAM}[\text{SP}+0]$
y	$=$	$\text{RAM}[\text{SP}+1]$
n	$=$	$\text{RAM}[\text{SP}+2]$
z	$=$	$\text{RAM}[\text{SP}+6]$

; continued

Examp

PowerSum_while1:

```
z = 0 ; y = 1 ; x = X ; n = N ;
while (n != 0)
{ while (n % 2 == 0)
  { y = y * (1 + x) ; x = x * x ; n = n / 2 ; }
  z = z + y ; y = y * x ; n = n - 1 ;
}
```

PowerSum_dol:

...

...

PowerSum_od1:

LOAD R0 [SP+PowerSum_z] ; R0 ← z

ADD R0 [SP+PowerSum_y] ; R0 ← z + y

STOR R0 [SP+PowerSum_z] ; z ← R0

LOAD R0 [SP+PowerSum_y] ; R0 ← y

LOAD R1 [SP+PowerSum_x] ; R1 ← x

MULS R0 R1 ; R0 ← y × x

STOR R0 [SP+PowerSum_y] ; y ← R0

LOAD R0 [SP+PowerSum_n] ; R0 ← n

SUB R0 1 ; R0 ← n - 1

STOR R0 [SP+PowerSum_n] ; n ← R0

BRA PowerSum_while0 ; jump back to beginning of

; outer repetition

x = RAM[SP+0]
y = RAM[SP+1]
n = RAM[SP+2]
z = RAM[SP+6]

Example: Translating a Java function into Assembly (8)

```
PowerSum_while1: LOAD
                  AND
                  BNE
```

```
PowerSum_do1:    ...
```

```
                BRA    PowerSum_while1    ; jump back to the inner loop
                                                ; inner repetition
```

```
PowerSum_od1 :   ...
                ...
```

```
                BRA    PowerSum_while0    ; jump back to the
                                                ; outer repetition
```

```
PowerSum_od0 :  ADD SP local_vars          ; remove local vars
                RTS                        ; result value is already
                                                ; there: just RETURN
```

```
x = RAM[SP+0]
y = RAM[SP+1]
n = RAM[SP+2]
z = RAM[SP+6]
```

```
{ y = y * (1+x) ; x = x * x ; n = n / 2 ; }
  z = z + y ; y = y * x ; n = n - 1 ;
}
return z ;
```

Questions?



Software interrupts or software exceptions.

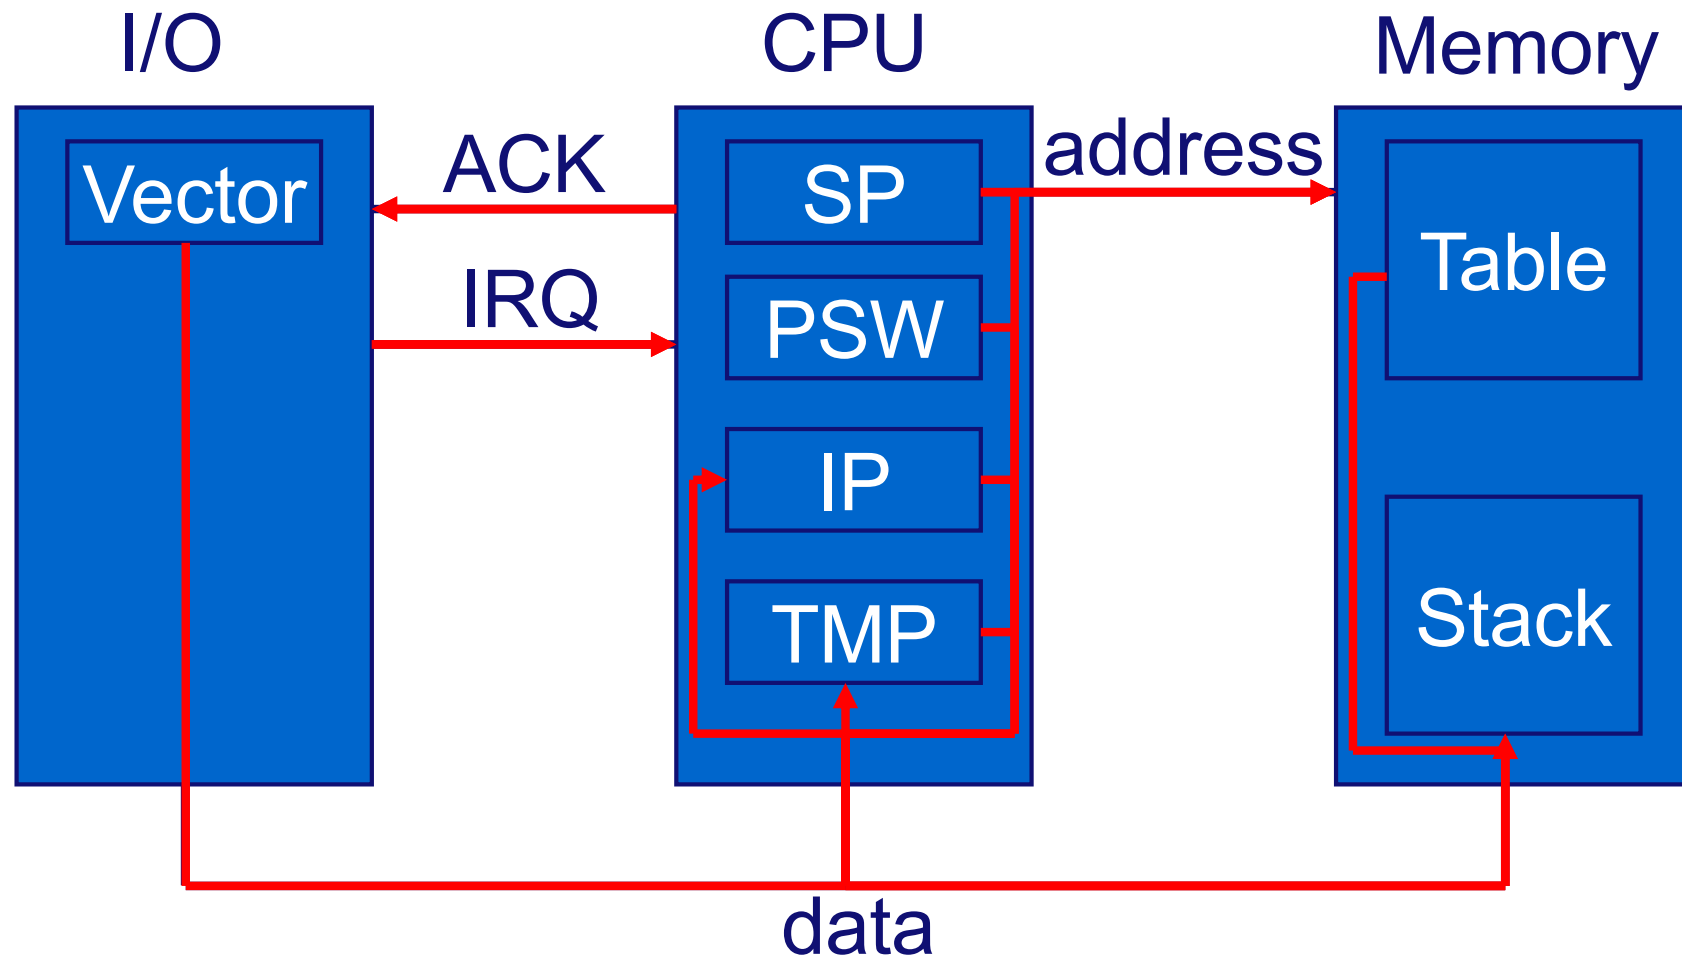
- ❑ SWI #number. ; ARM processor.
- ❑ TRAP ; PP2 processor.

- ❑ Call a special subroutine (in **supervisor mode on the ARM**).
- ❑ The effect of a software interrupt is:
 - LR := PC-4 (points to the next instruction)
 - PC := 8 (at address 8 there is a jump to the required software handling code)
 - CPSR := APSR (save the status flags).

Hardware interrupts happen

- ❑ Device passes interrupt to CPU
- ❑ CPU finishes current instruction and sends an acknowledgement to the device
- ❑ Device provides additional information
(e.g. IRQ-number, target address or table-index)
- ❑ {CPU fetches information and temporarily stores this}
- ❑ CPU pushes IP (and often PSW) onto the stack
- ❑ CPU calls interrupt handler
(which one depends on the additional information)

Hardware: interrupt (example)



Software: handle interrupt

- ❑ Save all registers.
- ❑ Find out which device caused an interrupt
- ❑ Get additional information (status etc).
- ❑ Handle the interrupt
(e.g. I/O error, start next program etc.)
- ❑ Signal to the device that the interrupt was serviced
- ❑ Restore all registers
- ❑ Execute **Re**Turn from **I**nterrupt instruction, or **ERET**.

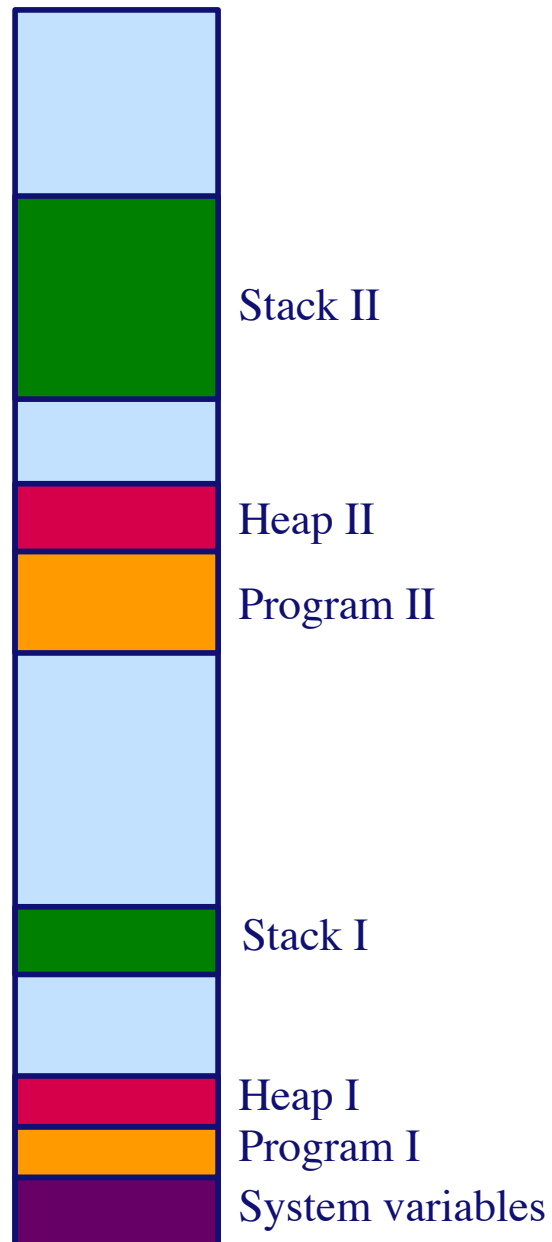
Interrupts: remarks

- ❑ Interrupt handlers must be transparent.
- ❑ If execution of an interrupt handler takes too long (longer than the period of the interrupt) the main program will not proceed.
- ❑ High priority interrupts are used for time critical applications (e.g. while burning a DVD, this was once modern... 😊)
- ❑ Interrupts are needed for multitasking.

Questions?

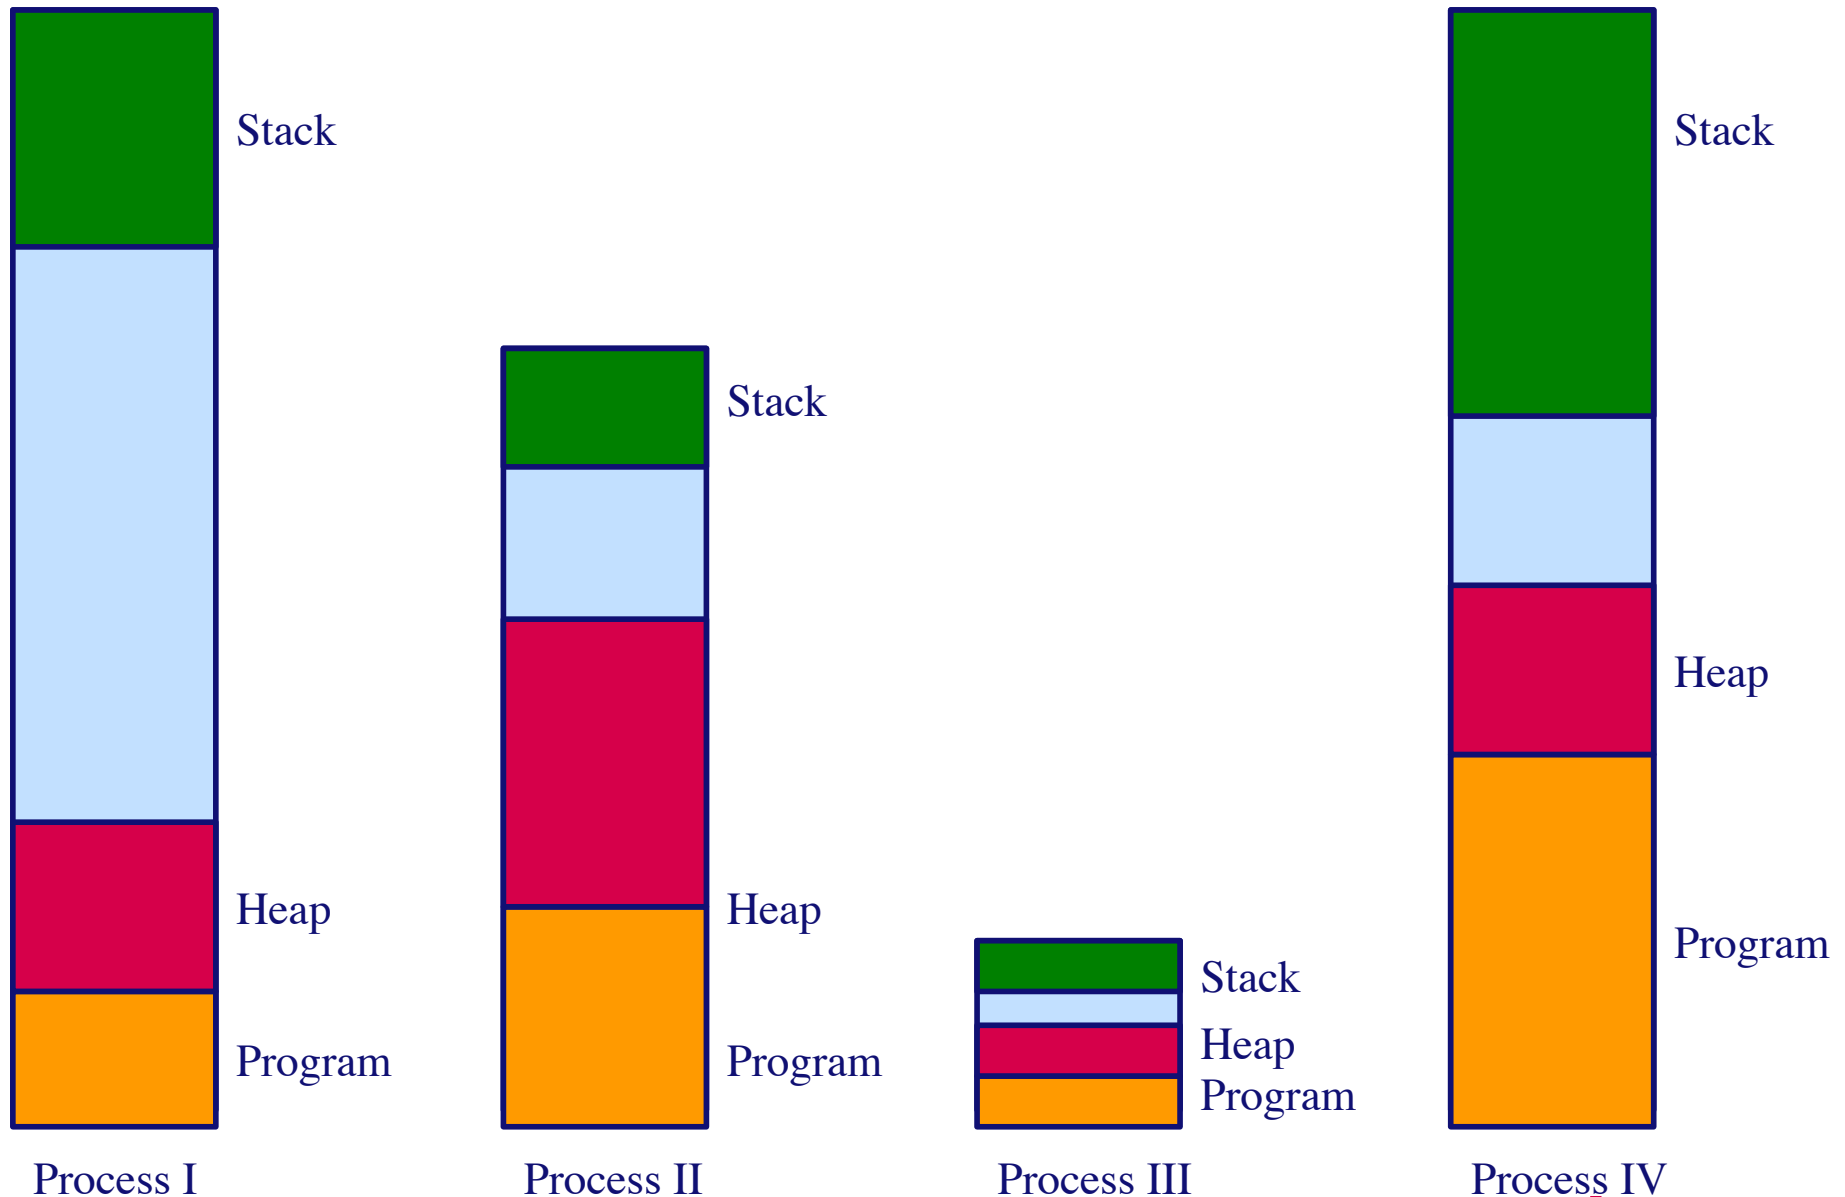


Elementary multitasking.

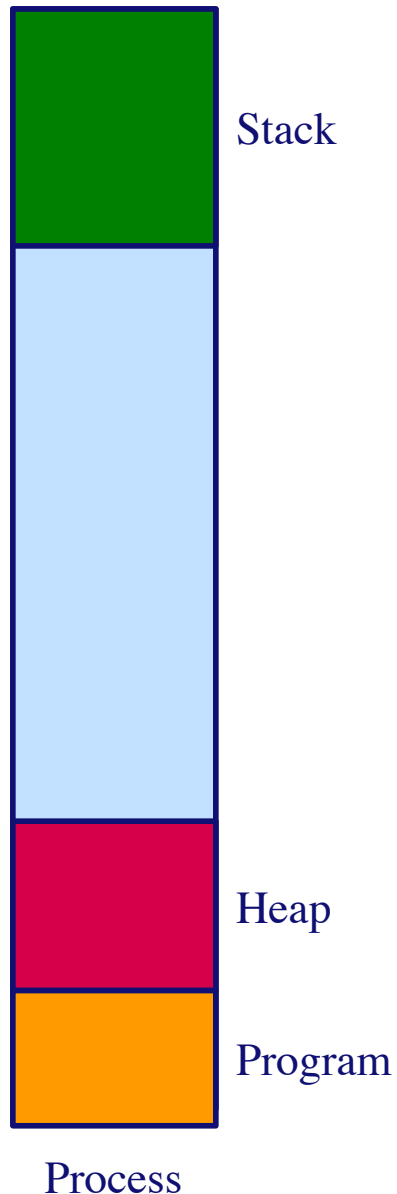


We can store more programs in memory, with their own stack and heap.

Another view on processes.



Each process has a process record



```
process_status (running, blocked, waiting,...)
```

```
stack_pointer
```

```
information_about_heap
```

```
information_about_access_rights
```

```
information_about_used_objects
```

```
etc. etc.
```

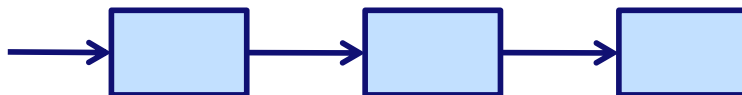
Running process



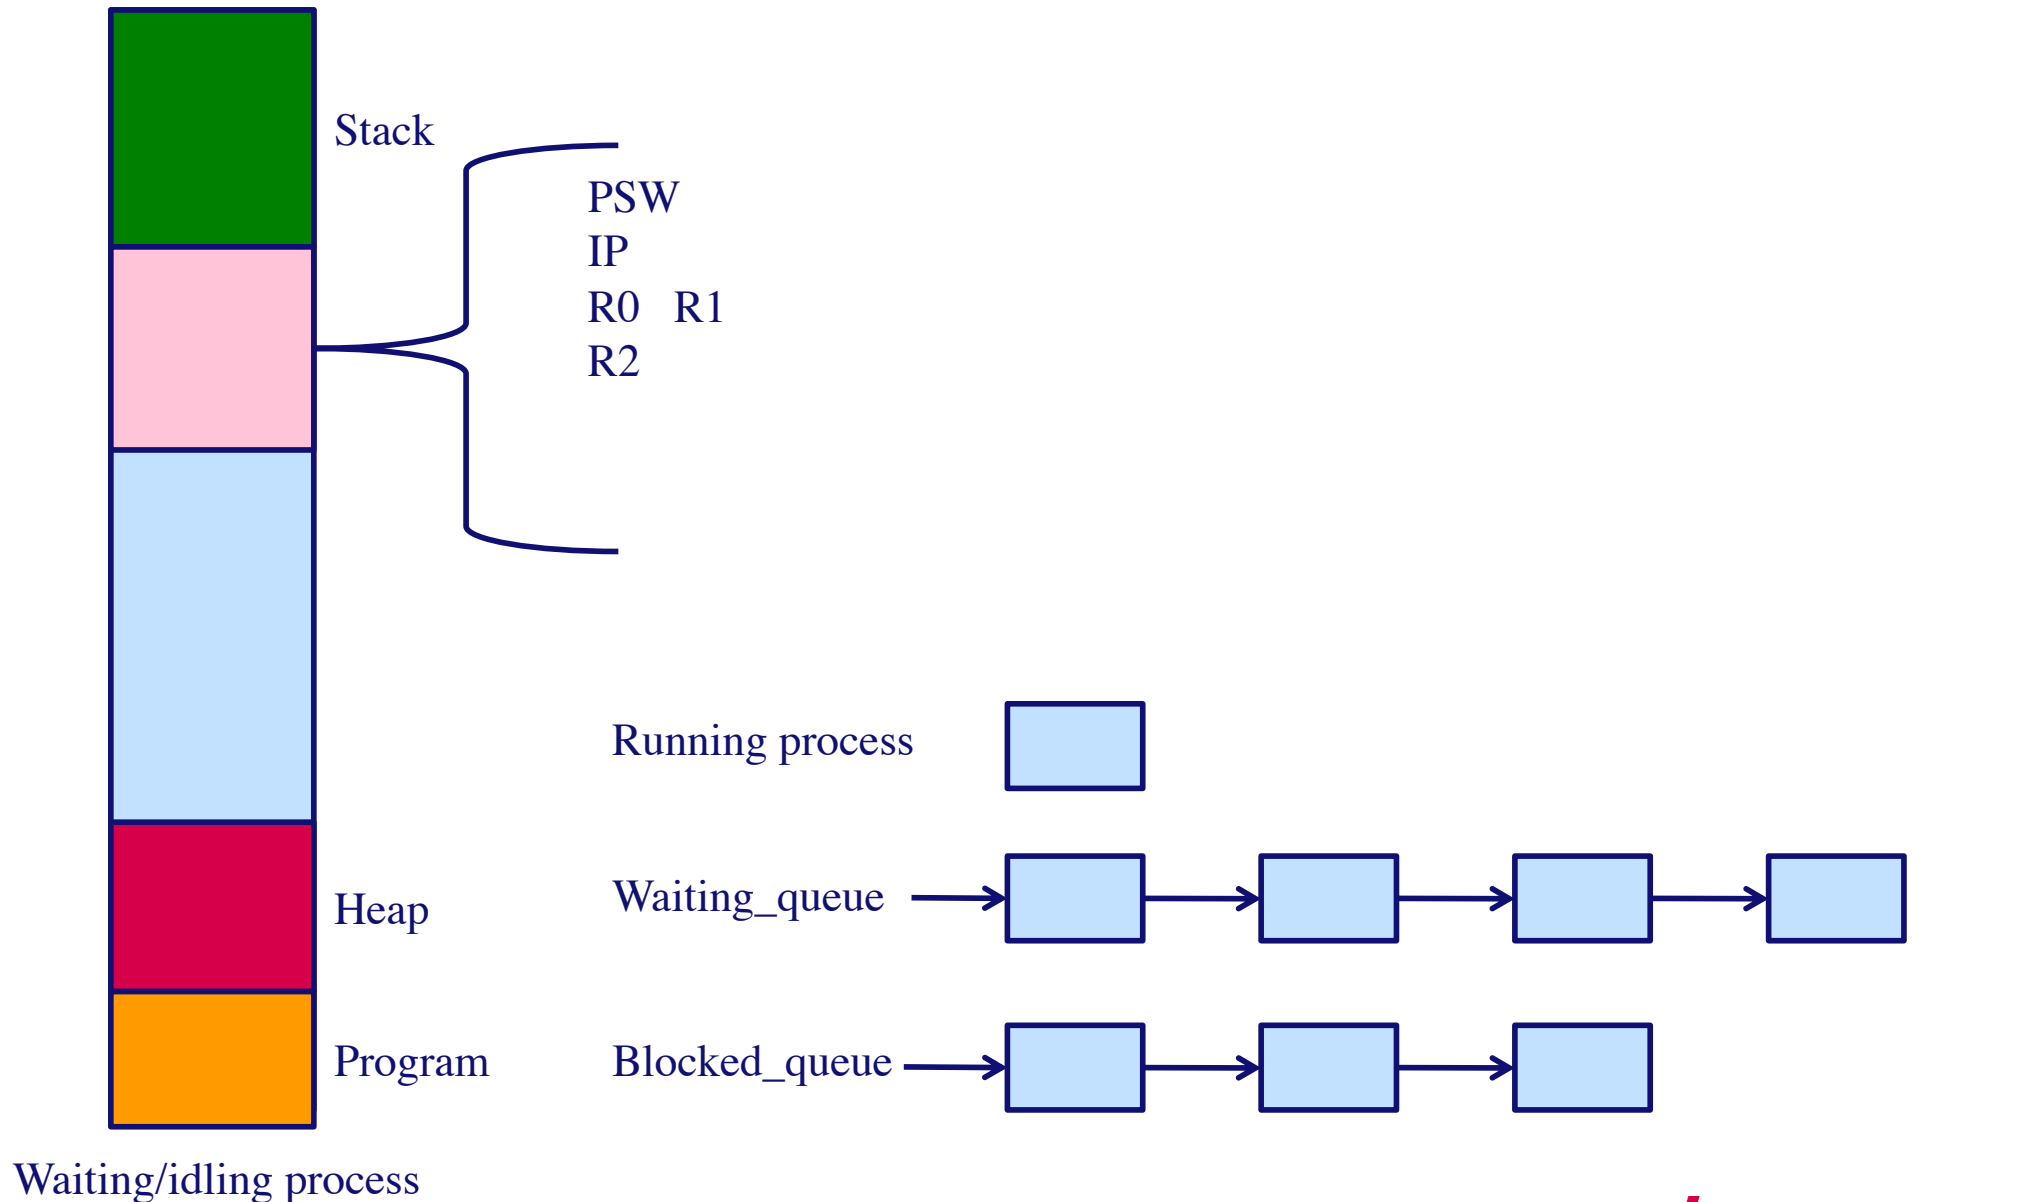
Waiting_queue



Blocked_queue



Blocked/waiting processes have info on the stack



Context switch

```
// When a process timer interrupt comes.
```

```
Save registers on stack of the current process;
```

```
If the waiting queue is not empty
```

```
{
```

```
  Save the stackpointer in the current process record;
```

```
  Set the status of the current process to wait;
```

```
  Put the process record at the end of the waiting list;
```

```
  Take a new process record from the waiting queue;
```

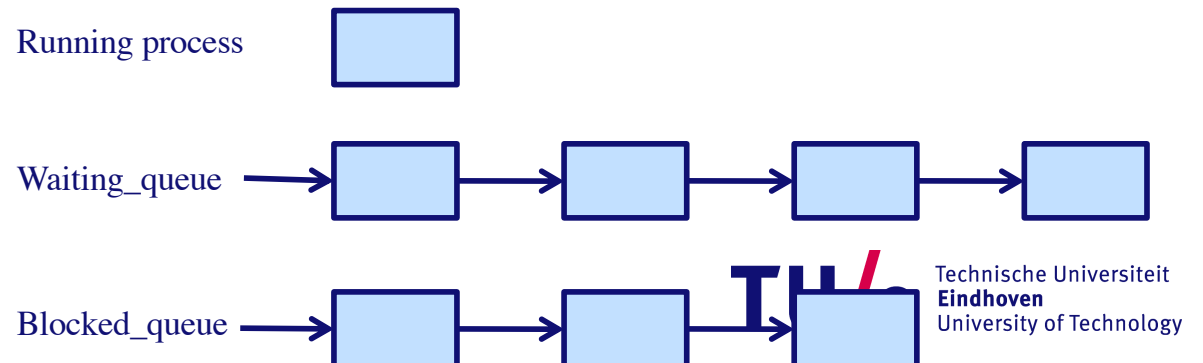
```
  Set its status to running;
```

```
  Put the stackpointer of the new process in the stack register.
```

```
}
```

```
Restore the registers of the new process;
```

```
RTI; Return from interrupt.
```



Context switch

```
// When a process timer interrupt comes.
```

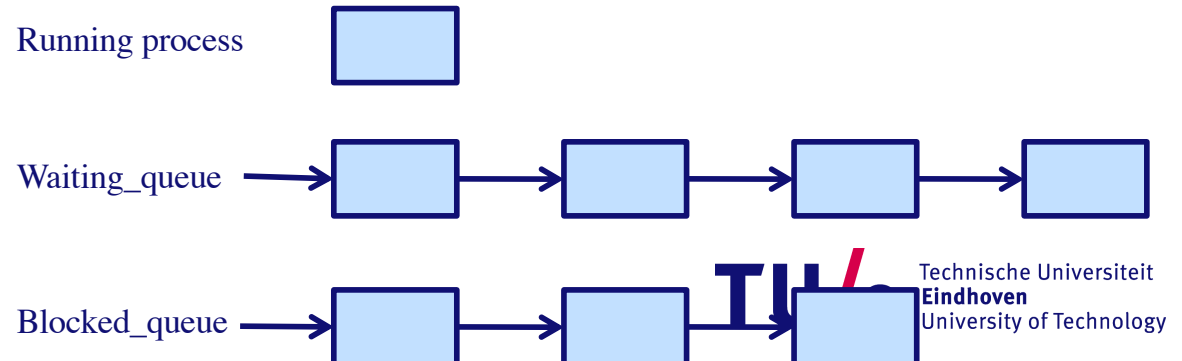
```
; Save registers on stack of the current process  
PUSH R0;
```

```
; If the waiting queue is not empty  
LOAD R0 [waiting_queue];  
BEQ finish_context_switch;
```

```
PUSH R1;  
PUSH R2;
```

```
; Save the stackpointer in the current process record;  
LOAD R1 [running_process]  
STOR SP [R1 + stackpointer];
```

```
; Set the status of current process to waiting;  
LOAD R2 waiting_status;  
STOR R2 [R1 + status];
```



Context switch

```
; Put the process record at the end of the waiting list;
```

```
    LOAD R2 [last_waiting_process];  
    STOR R1 [R2+next];  
    STOR R1 [last_waiting_process];
```

```
; Remove the first process record from the waiting queue;
```

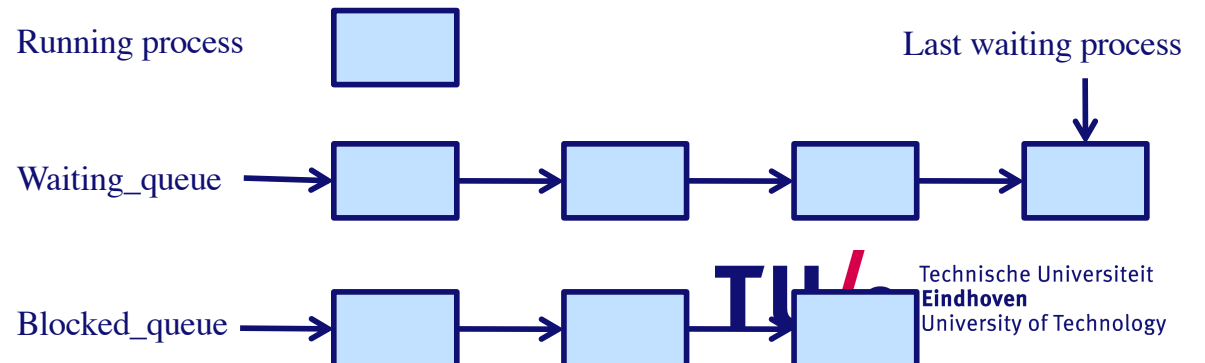
```
    ....
```

```
; Set its status to running;
```

```
    LOAD R1 running_status;  
    STOR R1 [R0 + status];           R0 points to the new process.
```

```
; Put the stackpointer of the new process in the stack register.
```

```
    LOAD SP [R0 + stackpointer]
```

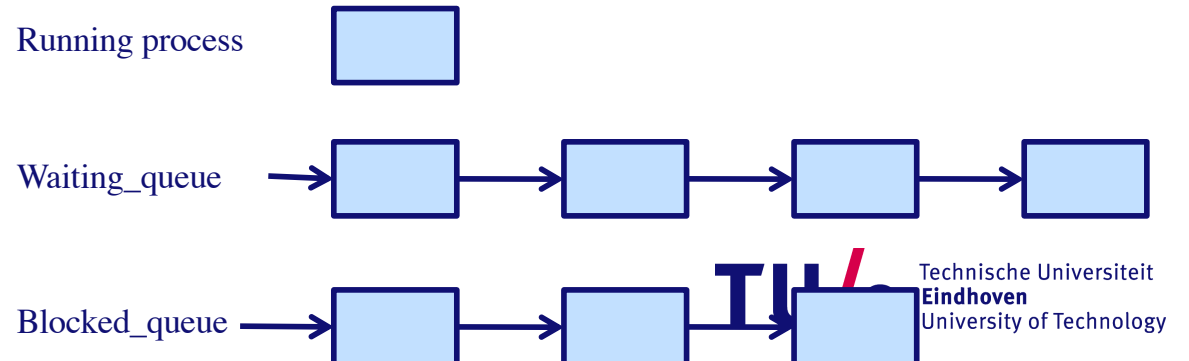


Context switch

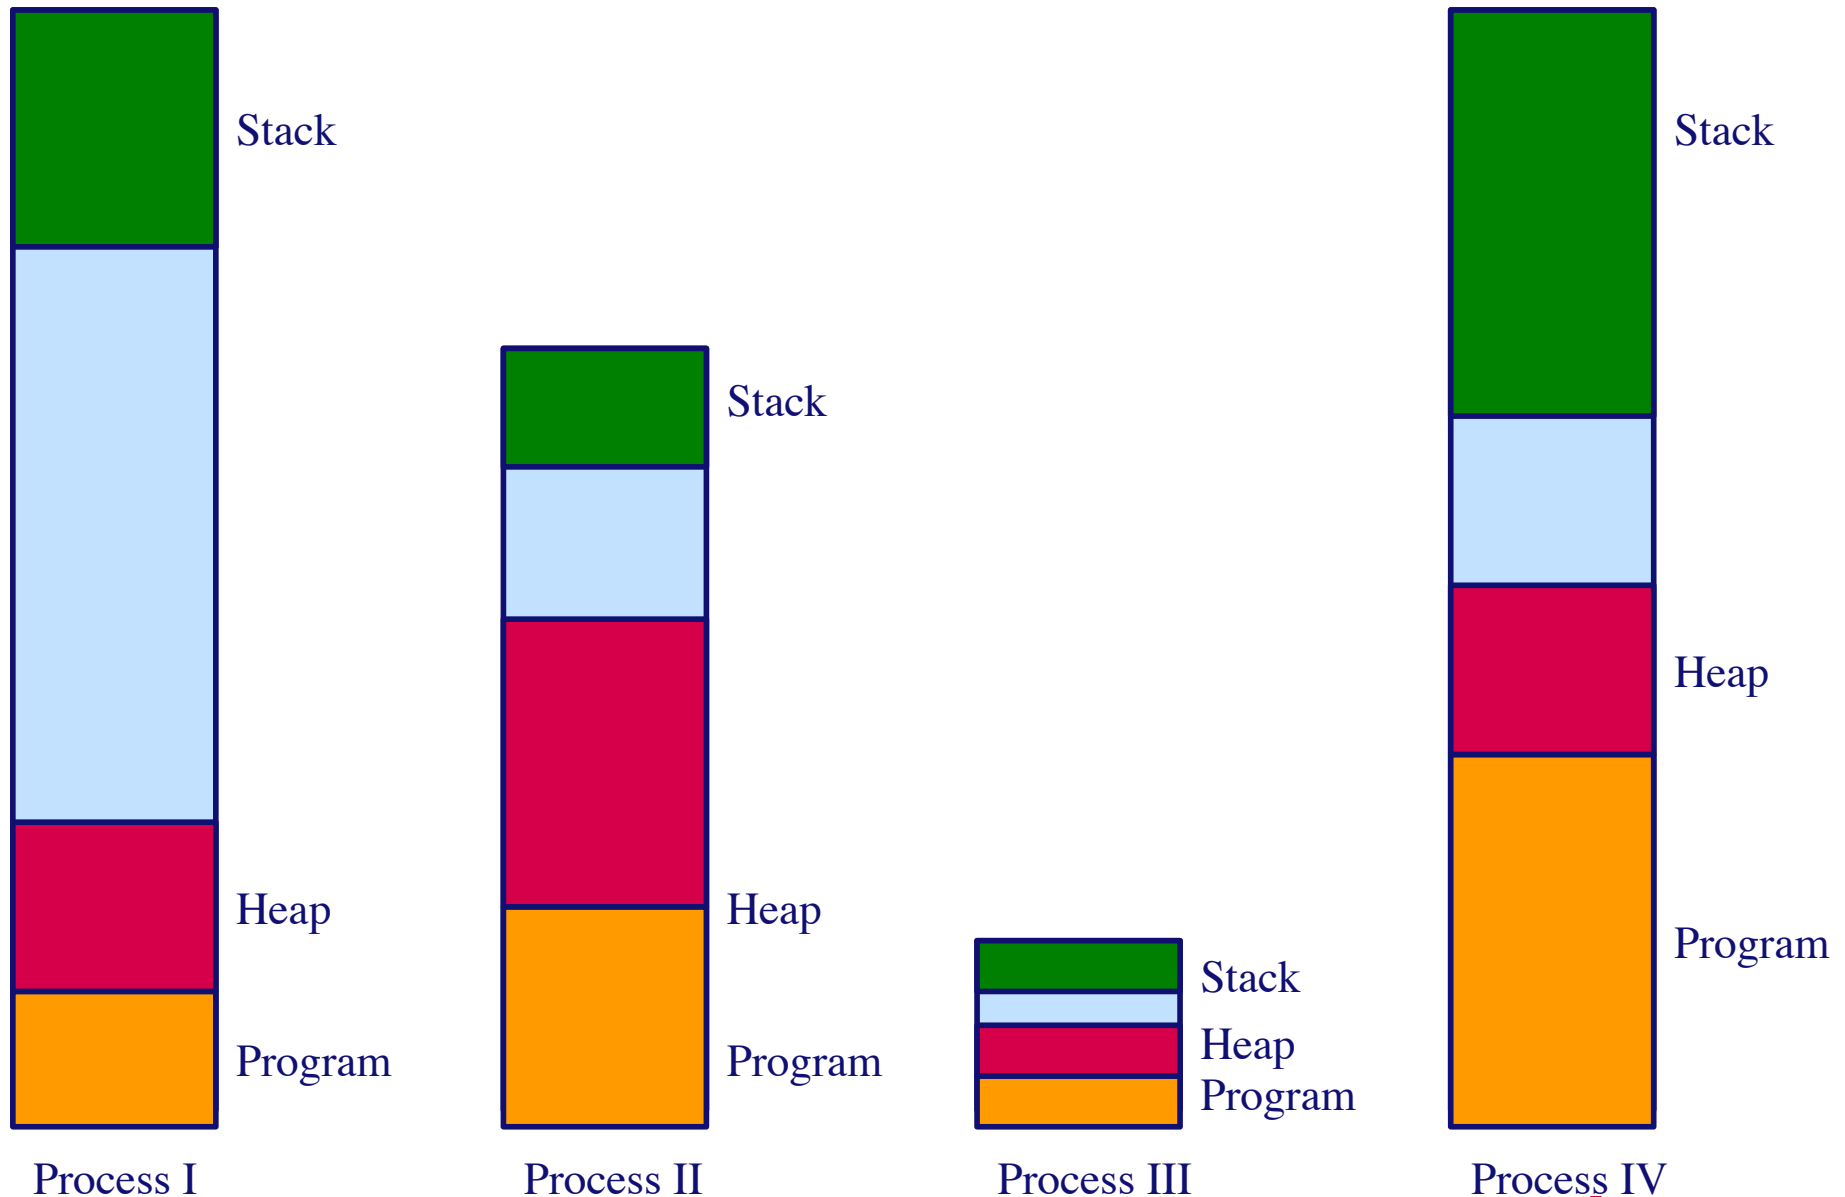
```
PULL R2;  
PULL R1;
```

```
finish_context_switch:  
; Restore the registers of the new process;  
PULL R0;
```

```
RTI; Return from interrupt.
```



The processor jumps from process to process.



Questions?



Non atomicity.

Translation to assembly code:

```
LOAD R0    [BP+x]  
ADD R0  1  
STOR R0    [BP+x]
```

Process 1:
 $x:=x+1$

Process 2:
 $x:=x+1$

$x=1$

Executing both processes can result in

$x=2$

when both processes are executed simultaneously.
Probability is low. This is hard to test.

Semaphores/mutexes: tool for process synchronisation



Passeer (Pass):
if passable
then indicate non passability
else
move to blocking queue
queue for this semaphore

Vrijgeven (Release):
if queue is empty
then indicate passable
otherwise move a process in the queue
to the processor waiting queue.



Edsger Wiebe Dijkstra (1930-2002)

These days called mutex variables.

Questions?



Summary

What did we learn:

- We can write programs in assembly (PP2, ARM).
- The stack is used for parameters of functions, local variables, the computation of expressions and returning results.
- We can systematically translate higher level programs into assembly code. This is what compilers do.
- We saw how to transform a single processor machine into a multithreading machine.
- We saw the need for mutual exclusion, and know that semaphores/mutexes can be used for this purpose.