

2IL50 Data Structures

2023-24 Q3

Lecture 7: Binary Search Trees

Announcements

Practice version of Segment 1 Interim Test on Ans

Remaining tests will be published after the results are known

Segment 2 Interim Test this Thursday

New room distribution, finalized tomorrow

No new rooms

Additional time students now in Neuron 0.242

Bring your own scrap paper – loose sheets!

No electronic devices, also not once you are done – a book?

External mice are fine, external keyboards not.

The keyboard is not your enemy, don't bash it loudly!

Respect your fellow students: stay seated and quiet until the test is over!

The test auto-submits in Ans after 45 (55) min

Dynamic Sets

Dynamic sets

Dynamic sets

Sets that can grow, shrink, or otherwise change over time.

Two types of operations:

- queries return information about the set
- modifying operations change the set

Common queries

Search, Minimum, Maximum, Successor, Predecessor

Common modifying operations

Insert, Delete

Dictionary

Dictionary

stores a set S of **elements**, each with an associated **key** (integer value)

Operations

Search(S, k): returns a pointer to an element x in S with $x.\text{key} = k$,
or *NIL* if such an element does not exist.

Insert(S, x): inserts element x into S , that is, $S \leftarrow S \cup \{x\}$

Delete(S, x): removes element x from S

Implementing a dictionary

| | Search | Insert | Delete |
|--------------|------------------|-------------|-------------|
| linked list | $\Theta(n)$ | $\Theta(1)$ | $\Theta(1)$ |
| sorted array | $\Theta(\log n)$ | $\Theta(n)$ | $\Theta(n)$ |
| hash table | $\Theta(1)$ | $\Theta(1)$ | $\Theta(1)$ |

Hash table

Running times are average times and assume independent uniform hashing and a large enough table (for example, of size $2n$)

Today Binary search trees

Binary Search Trees

Binary search trees

Binary search trees are an important data structure for dynamic sets.

They can be used as both a dictionary and a priority queue.

They accomplish many dynamic-set operations in $O(h)$ time, where h = height of tree.

Tree terminology

Binary tree: every node has 0, 1, or 2 children

Root: top node (no parent)

Leaf: node without children

Subtree rooted at node x : all nodes below and including x

Depth of node x : length of path from root to x

Depth of tree: max. depth over all nodes

Height of node x : length of longest path from x to leaf

Height of tree: height of root


Level: set of nodes with same depth

Family tree terminology

Left/right child

Parent

Grandparent ...



Binary search trees

root of T denoted by $T.\text{root}$


internal nodes have four fields:

key (and possible other satellite data)

left : points to left child

right : points to right child


p : points to parent. $T.\text{root}.p = \text{NIL}$



Binary search trees

Keys are stored only in internal nodes!

There are binary search trees which store keys only in the leaves ...



Binary search trees

A binary tree is


- a leaf or
- a root node x with a binary tree as its left and/or right child

Binary-search-tree property


- if y is in the left subtree of x , then $y.\text{key} \leq x.\text{key}$
- if y is in the right subtree of x , then $y.\text{key} \geq x.\text{key}$

Keys don't have to be unique ...


can use stricter property, take care when balancing



Binary search trees



height $h = 2$



height $h = 4$

Binary-search-tree property

- if y is in the left subtree of x , then $y.\text{key} \leq x.\text{key}$
- if y is in the right subtree of x , then $y.\text{key} \geq x.\text{key}$


Tree walks

Binary search trees are recursive structures

→ recursive algorithms often work well

TreeWalk(x)

- 1 **RecurseLeft()**
- 2 **RecurseRight()**
- 3 Do something



PreorderTreeWalk

- 1 Do something
- 2 **RecurseLeft()**
- 3 **RecurseRight()**

InorderTreeWalk

- 1 **RecurseLeft()**
- 2 Do something
- 3 **RecurseRight()**

PostorderTreeWalk


- 1 **RecurseLeft()**
- 2 **RecurseRight()**
- 3 Do something

Inorder tree walk

Example: print all keys in order using an **inorder tree walk**

InorderTreeWalk(x)

```
1 if  $x \neq NIL$   
2   InorderTreeWalk( $x$ .left)  
3   print  $x$ .key  
4   InorderTreeWalk( $x$ .right)
```



Correctness: follows by induction from the binary search tree property

Running time?

Intuitively, $O(n)$ time for a tree with n nodes, since we visit and print each node once.

Inorder tree walk

InorderTreeWalk(x)

```
1 if  $x \neq NIL$ 
2     InorderTreeWalk( $x$ .left)
3     print  $x$ .key
4     InorderTreeWalk( $x$ .right)
```

Theorem

If x is the root of an n -node subtree, then the call **InorderTreeWalk**(x) takes $\Theta(n)$ time.

Proof:

- $T(n)$ takes small, constant amount of time on empty subtree

$$T(0) = c \text{ for some positive constant } c$$

- for $n > 0$ assume that left subtree has k nodes, right subtree $n - k - 1$

$$T(n) = T(k) + T(n - k - 1) + d \text{ for some positive constant } d$$


use substitution method ... to show: $T(n) = (c + d)n + c$

Querying a binary search tree

TreeSearch(x, k)

```
1 if  $x == NIL$  or  $k = x.key$ : return  $x$ 
2 if  $k < x.key$ 
3     return TreeSearch( $x.left, k$ )
4 else
5     return TreeSearch( $x.right, k$ )
```

Binary-search-tree property




Initial call: **TreeSearch**($T.root, k$)

- **TreeSearch**($T.root, 4$)
- **TreeSearch**($T.root, 2$)

Running time:

$\Theta(\text{length of search path})$

worst case $\Theta(h)$




Querying a binary search tree – iteratively

TreeSearch(x, k)

```
1 if  $x == NIL$  or  $k = x.key$ : return  $x$ 
2 if  $k < x.key$ 
3     return TreeSearch( $x.left, k$ )
4 else
5     return TreeSearch( $x.right, k$ )
```

Binary-search-tree property



IterativeTreeSearch(x, k)

```
1 while  $x \neq NIL$  and  $k \neq x.key$ 
2     if  $k < x.key$ :  $x = x.left$ 
3     else:  $x = x.right$ 
4 return  $x$ 
```

the iterative version is more efficient on most computers

Minimum and maximum

Binary-search-tree property guarantees that

- the minimum key is located in the leftmost node
- the maximum key is located in the rightmost node



TreeMinimum(x)

```
1 while  $x.\text{left} \neq \text{NIL}$   
2    $x = x.\text{left}$   
3 return  $x$ 
```

TreeMaximum(x)

```
1 while  $x.\text{right} \neq \text{NIL}$   
2    $x = x.\text{right}$   
3 return  $x$ 
```

Binary-search-tree property



Minimum and maximum

Binary-search-tree property guarantees that

- the minimum key is located in the leftmost node
- the maximum key is located in the rightmost node


TreeMinimum(x)

```
1 while  $x.\text{left} \neq \text{NIL}$ 
2    $x = x.\text{left}$ 
3 return  $x$ 
```

TreeMaximum(x)

```
1 while  $x.\text{right} \neq \text{NIL}$ 
2    $x = x.\text{right}$ 
3 return  $x$ 
```

Binary-search-tree property



Running time?

Both procedures visit nodes on a downward path from the root

→ $O(h)$ time

Successor and predecessor

Assume that all keys are distinct

Successor of a node x :

node y such that $y.key$ is the smallest key $> x.key$

(if x has the largest key, then we say x 's successor is *NIL*)

We can find y based entirely on the tree structure, no key comparisons are necessary ...

Successor and predecessor

Successor of a node x

node y such that $y.key$ is the smallest key $> x.key$

Two cases:

1. x has a non-empty right subtree
→ x 's successor is the minimum in x 's right subtree
2. x has an empty right subtree
→ x 's successor y is the node of which x is the predecessor
(x is the maximum in y 's left subtree)

*as long as we move to the left up the tree (move up through right children),
we're visiting smaller keys ...*


Successor and predecessor

TreeSuccessor(x)

```
1 if  $x.\text{right} \neq \text{NIL}$ 
2     return TreeMinimum( $x.\text{right}$ )
3  $y = x.p$ 
4 while  $y \neq \text{NIL}$  and  $x = y.\text{right}$ 
5      $x = y$ 
6      $y = x.p$ 
7 return  $y$ 
```

TreePredecessor is symmetric

- Successor of 15?
- Successor of 6?
- Successor of 4?
- Predecessor of 6?




Successor and predecessor

TreeSuccessor(x)

```
1 if  $x.\text{right} \neq \text{NIL}$   
2     return TreeMinimum( $x.\text{right}$ )  
3  $y = x.p$   
4 while  $y \neq \text{NIL}$  and  $x = y.\text{right}$   
5      $x = y$   
6      $y = x.p$   
7 return  $y$ 
```

TreePredecessor is symmetric

Running time? $O(h)$



Insertion

TreeInsert(T, z)

```
1   $y = NIL$ 
2   $x = T.root$ 
3  while  $x \neq NIL$ 
4       $y = x$ 
5      if  $z.key < x.key$ :  $x = x.left$ 
6      else:  $x = x.right$ 
7   $z.p = y$ 
8  if  $y == NIL$ 
9       $T.root = z$ 
10 else
11     if  $z.key < y.key$ :  $y.left = z$ 
12     else:  $y.right = z$ 
```

to insert value v , insert node z with
 $z.key = v$, $z.left = NIL$, and $z.right = NIL$

traverse tree down to find correct position for z


Insertion

TreeInsert(T, z)

```
1  $y = NIL$ 
2  $x = T.root$ 
3 while  $x \neq NIL$ 
4      $y = x$ 
5     if  $z.key < x.key$ :  $x = x.left$ 
6     else:  $x = x.right$ 
7  $z.p = y$ 
8 if  $y == NIL$ 
9      $T.root = z$ 
10 else
11     if  $z.key < y.key$ :  $y.left = z$ 
12     else:  $y.right = z$ 
```

to insert value v , insert node z with
 $z.key = v$, $z.left = NIL$, and $z.right = NIL$

traverse tree down to find correct position for z



Insertion


TreeInsert(T, z)

```
1  $y = NIL$ 
2  $x = T.root$ 
3 while  $x \neq NIL$ 
4      $y = x$ 
5     if  $z.key < x.key$ :  $x = x.left$ 
6     else:  $x = x.right$ 
7  $z.p = y$ 
8 if  $y == NIL$ 
9      $T.root = z$ 
10 else
11     if  $z.key < y.key$ :  $y.left = z$ 
12     else:  $y.right = z$ 
```

Running time? $O(h)$

to insert value v , insert node z with
 $z.key = v$, $z.left = NIL$, and $z.right = NIL$

traverse tree down to find correct position for z




Deletion

We want to delete node z

TreeDelete has three cases:

z has no children

- delete z by having z 's parent point to *NIL*, instead of to z



Deletion

We want to delete node z


TreeDelete has three cases:

z has no children

- delete z by having z 's parent point to *NIL*, instead of to z

z has one child

- delete z by having z 's parent point to z 's child, instead of to z



Deletion

We want to delete node z

TreeDelete has three cases:

z has no children


- delete z by having z 's parent point to *NIL*, instead of to z

z has one child

- delete z by having z 's parent point to z 's child, instead of to z

z has two children

- z 's successor y has either no or one child
delete y from the tree and replace z 's key and satellite data with y 's



Deletion

We want to delete node z

TreeDelete has three cases:

z has no children


- delete z by having z 's parent point to *NIL*, instead of to z

z has one child

- delete z by having z 's parent point to z 's child, instead of to z

z has two children

- z 's successor y has either no or one child
delete y from the tree and replace z 's key and satellite data with y 's



Running time?

$O(h)$

Minimizing the running time

All operations can be executed in time proportional to the height h of the tree
(*instead of proportional to the number n of nodes in the tree*)

Worst case: $\Theta(n)$

Solution: guarantee small height (*balance the tree*) $\Theta(\log n)$

Balanced Search Trees

Balanced Search trees

There are many methods to balance a search tree.

by **weight**

for each node the **number of nodes** in the left and the right subtree is approximately equal

by **height**

for each node the **height** of the left and the right subtree is approximately equal

by **degree**

all leaves have the same depth, but the degree of the nodes differs
(*hence not a binary search tree*)

Weight-balanced search trees

BB[α]-tree

binary search tree where for each pair of nodes x, y , with y being a child of x we have

$$\alpha \leq \frac{\text{number of leaves in subtree rooted at } y}{\text{number of leaves in subtree rooted at } x} \leq 1 - \alpha$$

where α is a positive constant with $\alpha \leq 1/3$


For the height $h(n)$ it holds that $h(n) \leq h((1 - \alpha)n) + 1$

Master theorem: $h(n) = \Theta(\log n)$

Ideally: α as close as possible to $1/3$

But: $\alpha = 1/3$ gives too little flexibility for updates

α just smaller than $1/3$ works fine




Height-balanced search trees

AVL-tree

binary search tree where for each node

$$| \text{height left subtree} - \text{height right subtree} | \leq 1$$



Theorem

An AVL-tree with n nodes has height $\Theta(\log n)$.

Height-balanced search trees

Theorem

An AVL-tree with n nodes has height $\Theta(\log n)$.

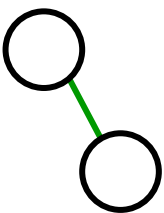
Proof

Let $n(h)$ = minimum number of nodes in an AVL-tree of height h

Claim: $n(h) \geq 2^{h/2}$

Proof of Claim: induction on h

$h = 0$  $n(0) = 1$ ✓

$h = 1$  $n(1) = 2$ ✓

Height-balanced search trees

Theorem

An AVL-tree with n nodes has height $\Theta(\log n)$.

Proof


Let $n(h)$ = minimum number of nodes in an AVL-tree of height h

Claim: $n(h) \geq 2^{h/2}$

Proof of Claim: induction on h

$$h \geq 2$$


$$\begin{aligned} n(h) &\geq 1 + n(h-1) + n(h-2) \\ &\geq 2n(h-2) \\ &\geq 2 \cdot 2^{(h-2)/2} \\ &= 2^{h/2} \end{aligned}$$



Degree-balanced trees

(2, 3)-tree

search tree where all leaves have the same depth and internal nodes have degree 2 or 3



Theorem

A (2, 3)-tree with n nodes has height $\Theta(\log n)$.

Red-black Trees

Another height-balanced search tree


Red-black trees

Red-black tree

binary search tree where each node has a color attribute which is either red or black

Red-black properties

1. Every node is either red or black.
2. The root is black.
3. Every leaf (*NIL*) is black.
4. If a node is red, then both its children are black.
(Hence no two reds in a row on a simple path from the root to a leaf)
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes.



Red-black trees: height

height of a node


number of edges on a longest path to a leaf

black-height of a node x


$bh(x)$ is the number of black nodes (including *NIL* leaves) on the path from x to leaf, not counting x .

Red-black properties

1. Every node is either red or black.
2. The root is black.
3. Every leaf (*NIL*) is black.
4. If a node is red, then both its children are black.
(Hence no two reds in a row on a simple path from the root to a leaf)
5. For each node, all paths from the node to descendant leaves contain the same number of black nodes.




Red-black trees: implementation detail



It is useful to replace each *NIL* by a single sentinel *T.nil* which is always black. The root's parent is also the sentinel.

Red-black trees: implementation detail



It is useful to replace each *NIL* by a single sentinel *T.nil* which is always black. The root's parent is also the sentinel.

NIL will not (always) be drawn on the following slides

Red-black trees

Lemma

A red-black tree with n nodes has height $\leq 2 \log(n + 1)$.

Proof

the subtree rooted at any node x contains
at least $2^{bh(x)} - 1$ internal nodes


the complete tree has n internal nodes

$$\rightarrow 2^{bh(T.\text{root})} - 1 \leq n$$

$$\rightarrow bh(T.\text{root}) \leq \log(n + 1)$$

$$\begin{aligned} \text{height}(T) &= \text{number of edges on a longest path to a leaf} \\ &\leq 2 \cdot bh(T.\text{root}) \\ &\leq 2 \log(n + 1) \end{aligned}$$

“smallest” subtree with
black-height $bh(x)$



Balanced binary search trees

Advantages of balanced binary search trees

over linked lists

efficient search in $\Theta(\log n)$ time

over sorted arrays

efficient insertion and deletion in $\Theta(\log n)$ time

over hash tables

can find successor and predecessor efficiently in $\Theta(\log n)$ time